

MTH321
Numerical Analysis Of Partial Differential Equations

Numerical Analysis of the Robertson Stiff Chemical Kinetics System

Rengerile Na 2359838

Suning Peng 2361786

Kainan Zhao 2362073

Yixun Wu 2144703

Keji Chen 2363225



Xi'an Jiaotong-Liverpool University
School of Mathematics and Physics

September 16, 2026

Contents

Contents	1
1 Introduction	2
1.1 Background and Motivation	2
1.2 Problem Statement and Report Roadmap	2
2 Mathematical Formulation	2
2.1 Reaction Kinetics Model	2
2.2 Conservation Invariant and Structural Zero Eigenvalue	3
2.3 Reduced Jacobian and Stiffness Ratio Definition	4
3 Numerical Integration Methods	6
3.1 Explicit Euler Method	6
3.2 Explicit Fourth-Order Runge–Kutta Method (RK4)	7
3.3 Implicit Euler Method with Newton–Raphson Iteration	9
3.4 Trapezoidal Method (Crank–Nicolson) as a Second Implicit Scheme	11
3.5 Adaptive Step-Size Control	12
4 Linear Stability Analysis	14
4.1 Absolute Stability Regions of the Four Methods	14
4.2 Eigenvalue Analysis of the Robertson System	15
4.3 Step-Size Constraints: Theory versus Experiment	17
5 Numerical Experiments and Validation	19
5.1 Reference Solution Construction	19
5.2 Sanity Checks	20
5.3 Convergence Order Verification	20
5.4 Explicit Euler: Stability Boundary and Physical Admissibility	22
5.5 Solver Diagnostics	23
5.5.1 Newton-Tolerance Sensitivity and Initial Guess	23
5.5.2 Adaptive Step-Size Control	24
5.6 Efficiency Comparison	24
6 Discussion	26
6.1 Physical Interpretation via Quasi-Steady-State Approximation (QSSA)	26
6.2 Why Explicit Methods Fail and Implicit Methods Succeed	30
6.3 Three-Fold Validation Strategy Recap	31
6.4 Limitations	31
7 Conclusions	32
8 References	33
9 Appendix	35

1 Introduction

1.1 Background and Motivation

Stiff initial value problems arise wherever processes with very different rates are coupled, most prominently in chemical kinetics. Their difficulty is not accuracy but stability: since Curtiss and Hirschfelder [1], it has been known that an explicit method may be forced to use a step size determined by the fastest transient, even after that transient has largely decayed. The reaction system introduced by Robertson [2] has become a standard benchmark for stiff ODE solvers. Its three species react through steps whose rate constants differ by a factor $k_2/k_1 = 7.5 \times 10^8$, so that the intermediate species is produced slowly but consumed almost instantly.

Beyond being a benchmark, the problem is a good test case for *validation*, for three reasons. First, no closed-form solution is available on the interval of interest, so accuracy must be established against an independently verified numerical reference. Second, the exact solution obeys a linear invariant and must stay non-negative, which gives diagnostics that are independent of accuracy and can fail independently of it. Third, the step size at which explicit methods break down can be predicted from the Jacobian, which allows a genuine prediction–observation comparison. We therefore treat the problem as a consultant would: the aim is not to reproduce a textbook example, but to produce a validated numerical answer and to show, quantitatively, why it can be trusted.

1.2 Problem Statement and Report Roadmap

We study the initial value problem for the concentrations of the three species, the coefficients of Eq. (4) follow the mass-action convention fixed by the problem statement, without an extra factor of two in the quadratic term. Since $\mathbf{1}^T \mathbf{f} \equiv 0$, the exact solution stays on the plane $y_1 + y_2 + y_3 = 1$ and is non-negative.

The report answers three questions. (i) Are the numerical solutions correct, i.e. do the implementations reproduce their formal orders against a verified reference? (ii) Which step-size restriction is active for each method, accuracy or stability, and does the solution remain physically admissible? (iii) What does each method cost at comparable accuracy? To answer them we implement Explicit Euler, RK4, Implicit Euler and Trapezoidal/Crank–Nicolson (with Newton’s method) and a step-doubling adaptive Euler method, derive their stability properties, and validate them against a tight Radau reference. In short, we find that stiffness appears here as a *stability* restriction, $h < 5.895 \times 10^{-4}$ for Explicit Euler, and that the implicit advantage depends on the accuracy demanded.


The report is organised as follows. Section 2 presents the model, its invariant and its Jacobian. Section 3 describes the numerical methods, including the Newton iteration and the adaptive controller. Section 4 derives the stability regions, the eigenvalue structure and the step-size constraints. Section 5 reports the experiments: reference construction, convergence orders, stability and admissibility, solver diagnostics and efficiency. Section 6 interprets the results, through a quasi-steady-state analysis (Section 6.1), the explicit-versus-implicit comparison (Section 6.2), a recap of the validation strategy (Section 6.3) and the limitations (Section 6.4). Section 7 concludes.

2 Mathematical Formulation

2.1 Reaction Kinetics Model

We consider the classical stiff chemical kinetics problem due to Robertson [2], which models an autocatalytic reaction network among three chemical species A , B , and C . The scheme consists

of three elementary steps operating on widely separated time scales:



with rate constants spanning approximately eight orders of magnitude (k_1 to k_2). This disparity in reaction rates is the source of the stiffness analyzed in Section 4.

The reaction labels above are mnemonic descriptions of the underlying chemistry; the governing ordinary differential equation system is the authoritative model and fixes the mass-action coefficient convention used throughout this report. With $y = (y_1, y_2, y_3)^\top$ denoting the concentrations of A , B , and C respectively,

$$\begin{aligned} y_1' &= -0.04 y_1 + 10^4 y_2 y_3, \\ y_2' &= 0.04 y_1 - 10^4 y_2 y_3 - 3 \times 10^7 y_2^2, \\ y_3' &= 3 \times 10^7 y_2^2, \end{aligned} \quad (4)$$

subject to the initial condition $y(0) = (1, 0, 0)^\top$.

Each term admits a direct kinetic interpretation: the $-0.04 y_1$ / $+0.04 y_1$ pair reflects the depletion of A and production of B by reaction 1; the $10^4 y_2 y_3$ pair reflects the regeneration of A from B at the expense of B via reaction 3 (with C acting as a catalyst, left unchanged); and the $3 \times 10^7 y_2^2$ terms reflect the self-reaction of B via reaction 2, which both consumes B and produces C .

2.2 Conservation Invariant and Structural Zero Eigenvalue

The mathematical foundation of the Robertson chemical kinetics model rests on an exact linear mass conservation law and an intrinsic structural degeneracy of its governing vector field. Prior to performing any numerical integration or dynamic reduction, we establish these fundamental algebraic properties formally.

Proposition 1 (Mass Conservation and Structural Zero Eigenvalue). *Let $f(y) = [f_1(y), f_2(y), f_3(y)]^\top : \mathbb{R}^3 \rightarrow \mathbb{R}^3$ be the vector field of the Robertson reaction system defined by*

$$f_1(y) = -k_1 y_1 + k_3 y_2 y_3, \quad (5a)$$

$$f_2(y) = k_1 y_1 - k_3 y_2 y_3 - k_2 y_2^2, \quad (5b)$$

$$f_3(y) = k_2 y_2^2, \quad (5c)$$

where $k_1 = 0.04$, $k_2 = 3 \times 10^7$, and $k_3 = 10^4$. Let $\mathbf{1} = [1, 1, 1]^\top \in \mathbb{R}^3$ denote the ones vector. Then:

1. **Algebraic Identity:** The vector field satisfies $\mathbf{1}^\top f(y) = 0$ identically for all state vectors $y \in \mathbb{R}^3$.
2. **Exact Conservation Law:** Any exact trajectory $y(t)$ subject to initial condition $y(0) = [1, 0, 0]^\top$ remains confined to the two-dimensional invariant simplex hyper-plane

$$\mathcal{M} = \{y \in \mathbb{R}^3 \mid y_1 + y_2 + y_3 = 1\} \quad \forall t \geq 0. \quad (6)$$

3. **Structural Zero Eigenvalue:** The full 3×3 Jacobian matrix $J_f(y) \equiv \frac{\partial f}{\partial y}(y)$ possesses a zero eigenvalue $\lambda_3 = 0$ with associated constant left eigenvector $\mathbf{1}^\top$ everywhere in state space, independent of y or time t .

Proof. We prove each statement sequentially:

1. Algebraic Identity: For an arbitrary vector $y = [y_1, y_2, y_3]^T \in \mathbb{R}^3$, expanding the inner product $\mathbf{1}^T f(y) = f_1(y) + f_2(y) + f_3(y)$ yields

$$\mathbf{1}^T f(y) = (-k_1 y_1 + k_3 y_2 y_3) + (k_1 y_1 - k_3 y_2 y_3 - k_2 y_2^2) + (k_2 y_2^2). \quad (7)$$

Pairwise cancellation of identical reaction terms gives:

$$\mathbf{1}^T f(y) = \underbrace{(-k_1 y_1 + k_1 y_1)}_{=0} + \underbrace{(k_3 y_2 y_3 - k_3 y_2 y_3)}_{=0} + \underbrace{(-k_2 y_2^2 + k_2 y_2^2)}_{=0} = 0. \quad (8)$$

Because this cancellation depends solely on the stoichiometric coefficients of the reaction network and holds without assuming y satisfies the differential equations or non-negativity constraints, $\mathbf{1}^T f(y) \equiv 0$ is a global algebraic identity over $\mathbb{R}^3$.

2. Exact Conservation Law: Along any solution trajectory $y(t)$ governed by $\frac{dy}{dt} = f(y(t))$, the time derivative of the total chemical species sum is

$$\frac{d}{dt} (y_1(t) + y_2(t) + y_3(t)) = \frac{d}{dt} (\mathbf{1}^T y(t)) = \mathbf{1}^T \frac{dy}{dt} = \mathbf{1}^T f(y(t)) \equiv 0. \quad (9)$$

Integrating from $t = 0$ to t gives $\mathbf{1}^T y(t) = \mathbf{1}^T y(0)$. Substituting the benchmark initial condition $y(0) = [1, 0, 0]^T$ yields $\mathbf{1}^T y(0) = 1 + 0 + 0 = 1$. Hence, $y_1(t) + y_2(t) + y_3(t) \equiv 1$ holds for all $t \geq 0$.

3. Structural Zero Eigenvalue: Differentiating the algebraic identity $\mathbf{1}^T f(y) = 0$ with respect to the state vector y via the multivariate chain rule gives

$$\frac{\partial}{\partial y} (\mathbf{1}^T f(y)) = \mathbf{1}^T \frac{\partial f}{\partial y}(y) = \mathbf{1}^T J_f(y) = \mathbf{0}^T. \quad (10)$$

Transposing Eq. (10) yields $J_f(y)^T \mathbf{1} = \mathbf{0} = \mathbf{0} \cdot \mathbf{1}$. Because $\mathbf{1} \neq \mathbf{0}$, this confirms that $\lambda = 0$ is an exact eigenvalue of $J_f(y)^T$ (and consequently of $J_f(y)$) for all $y \in \mathbb{R}^3$. Crucially, this zero eigenvalue is *structural*: it originates entirely from mass conservation embedded in the network stoichiometry rather than local state-dependent numerical coincidences. $\square$

Implications for Stiffness Diagnostic Metrics The existence of a persistent structural zero eigenvalue directly impacts standard numerical stiffness diagnostics. The conventional stiffness ratio $S(t)$ for an ODE system is defined by the spread of Jacobian eigenvalues:

$$S_{\text{naive}}(t) = \frac{\max_j |\operatorname{Re}(\lambda_j(t))|}{\min_j |\operatorname{Re}(\lambda_j(t))|}. \quad (11)$$

If evaluated directly on the full 3×3 Jacobian matrix $J_f(y)$, Proposition 1 dictates that $\min_j |\operatorname{Re}(\lambda_j(t))| \equiv \lambda_3 = 0$ at all times. Substituting this into Eq. (11) results in a division-by-zero defect, rendering $S_{\text{naive}}(t)$ ill-defined or infinitely large regardless of the actual dynamic stiffness. This mathematical singularity underscores the necessity of eliminating the redundant degree of freedom prior to diagnostic evaluation, leading directly to the state-space reduction method formulated in Section 2.3.

2.3 Reduced Jacobian and Stiffness Ratio Definition

Proposition before shows that the exact trajectory is confined to the invariant manifold $M = \{y \in \mathbb{R}^3 : y_1 + y_2 + y_3 = 1\}$ and that the full 3×3 Jacobian carries a structural zero eigenvalue. This section makes that reduction explicit: we eliminate y_3 , derive the resulting 2×2 Jacobian J_r , and prove that its eigenvalues coincide exactly with the two non-trivial eigenvalues of J_f . This J_r is then the object evaluated numerically in Section 2.2 and reused by the Newton solver in Sections 3.3–3.4.

Reduced two-dimensional system. Substituting $y_3 = 1 - y_1 - y_2$ into f_1 and f_2 from Eqs. 5a-5b defines a self-contained system on $\mathbb{R}^2$,

$$\dot{y}_1 = g_1(y_1, y_2), \quad \dot{y}_2 = g_2(y_1, y_2), \quad (12)$$

with

$$g_1(y_1, y_2) = -k_1 y_1 + k_3 y_2 (1 - y_1 - y_2), \quad (13)$$

$$g_2(y_1, y_2) = k_1 y_1 - k_3 y_2 (1 - y_1 - y_2) - k_2 y_2^2. \quad (14)$$

Since $g_1(y_1, y_2) = f_1(y_1, y_2, 1 - y_1 - y_2)$ and likewise for g_2 , this is not a separate model: it is f restricted to M and expressed in the affine chart $\iota(y_1, y_2) = (y_1, y_2, 1 - y_1 - y_2)$, which is a smooth bijection from $\mathbb{R}^2$ onto M with smooth inverse $\pi(y_1, y_2, y_3) = (y_1, y_2)$.

Proposition 2 (Reduced Jacobian). *Writing $y_3 = 1 - y_1 - y_2$, the Jacobian of the reduced system (12) is*

$$J_r(y_1, y_2) = \begin{pmatrix} -k_1 - k_3 y_2 & k_3(y_3 - y_2) \\ k_1 + k_3 y_2 & -k_3(y_3 - y_2) - 2k_2 y_2 \end{pmatrix}. \quad (15)$$

Its trace and determinant reduce to

$$\text{tr } J_r = -k_1 - k_3 y_3 - 2k_2 y_2, \quad \det J_r = 2k_2 y_2 (k_1 + k_3 y_2), \quad (16)$$

using the identity $(J_r)_{11} = -(J_r)_{21}$, which itself follows from $g_1 + g_2 = -f_3(y_1, y_2, y_3) = -k_2 y_2^2$ being independent of y_1 .

Proof. Direct differentiation of g_1, g_2 with respect to y_1, y_2 , or equivalently the chain rule $Dg = D\pi \cdot J_f(\iota(y_1, y_2)) \cdot D\iota$ with $D\iota = \begin{pmatrix} 1 & 0 \\ 0 & 1 \\ -1 & -1 \end{pmatrix}$, gives Eq. (15). The trace/determinant simplification follows by direct expansion and use of $(J_r)_{11} + (J_r)_{21} = 0$. $\square$

Exact reduction, not approximate elimination. It remains to show that the two eigenvalues of J_r are exactly the two non-trivial eigenvalues of the full 3×3 Jacobian J_f — i.e. that eliminating y_3 is an exact restriction to the invariant tangent space, not a heuristic dimension reduction.

Proposition 3 (Exact Correspondence of Spectra). *Let $T = \{v \in \mathbb{R}^3 : \mathbf{1}^T v = 0\}$ be the tangent space of M , and let $\Phi : T \rightarrow \mathbb{R}^2$, $\Phi(v_1, v_2, v_3) = (v_1, v_2)$, be the linear isomorphism induced by π . Then for every $y \in \mathbb{R}^3$,*

$$\text{spec}(J_f(y)) = \{0\} \cup \text{spec}(J_r(\pi(y))), \quad (17)$$

where J_r is evaluated at $\pi(y) = (y_1, y_2)$ using Eq. (15) with $y_3 = 1 - y_1 - y_2$.

Proof. Differentiating the algebraic identity $\mathbf{1}^T f(y) \equiv 0$ (Proposition 1, part 1) with respect to y gives $\mathbf{1}^T J_f(y) \equiv 0$ for every $y \in \mathbb{R}^3$; that is, $\mathbf{1}$ is a left null vector of $J_f(y)$, equivalently $\text{im } J_f(y) \subseteq T$ for every input, not only for inputs already in T . Choosing a basis of $\mathbb{R}^3$ adapted to the splitting $\mathbb{R}^3 = T \oplus \text{span}\{e\}$ (any $e \notin T$), the image condition forces $J_f(y)$ into the block-triangular form

$$J_f(y) \sim \begin{pmatrix} A & b \\ 0 & 0 \end{pmatrix},$$

where A is the 2×2 matrix representing $J_f(y)|_T : T \rightarrow T$ in that basis. Hence $\text{spec}(J_f(y)) = \{0\} \cup \text{spec}(A)$, recovering the structural zero eigenvalue of Proposition 1, part 3, together with the two remaining eigenvalues carried by A .

It remains to identify A with J_r . Because $\text{im } J_f(y) \subseteq T$ in particular for inputs drawn from T , the restriction $J_f(y)|_T : T \rightarrow T$ is a well-defined linear operator, and since Φ is a linear isomorphism $T \rightarrow \mathbb{R}^2$, A is similar to $\Phi \circ J_f(y)|_T \circ \Phi^{-1}$ regardless of which basis of T was used above. Direct computation shows $\Phi \circ J_f(y)|_T \circ \Phi^{-1} = D\pi \cdot J_f(y) \cdot D\iota = J_r(\pi(y))$, i.e. exactly the matrix of Proposition 2. Similar matrices share eigenvalues, so $\text{spec}(A) = \text{spec}(J_r(\pi(y)))$, which gives Eq. (17). $\square$

Proposition 3 justifies working with J_r throughout the remainder of the report: its two eigenvalues are the two non-trivial eigenvalues of the full system, exactly, at every point on M — not an approximation obtained by discarding a row and column.

Stiffness ratio.

Definition 1 (Stiffness Ratio). *Let $\lambda_1(t), \lambda_2(t)$ denote the eigenvalues of $J_r(y_1(t), y_2(t))$ along a trajectory. Provided both eigenvalues satisfy $|\lambda_j(t)| > \varepsilon_\lambda$ for a fixed numerical threshold ε_λ separating a genuinely nonzero eigenvalue from floating-point noise, define*

$$S(t) = \frac{\max_j |\text{Re } \lambda_j(t)|}{\min_j |\text{Re } \lambda_j(t)|}. \quad (18)$$

The threshold ε_λ is needed because, from Eq. (16), $\det J_r = 2k_2y_2(k_1 + k_3y_2)$ vanishes exactly whenever $y_2 = 0$ — in particular at the initial condition $y_2(0) = 0$, where one eigenvalue of J_r is exactly zero and $S(t)$ is undefined. This is not a numerical artifact but an exact algebraic degeneracy at $t = 0$; accordingly, $S(t)$ is only plotted from the first positive output time $t = 10^{-8}$ onward, and ε_λ is calibrated in Section 4.2 once $\lambda_1(t), \lambda_2(t)$ have been evaluated numerically along the reference trajectory.

Interface to later sections. The closed-form J_r derived here serves two roles beyond the present section: (i) Section 4.2 evaluates $\lambda_{1,2}(t)$ and $S(t)$ from Eq. (16) along the reference trajectory of Section 5.1; and (ii) Newton’s method for implicit Euler and trapezoidal (Sections 3.3–3.4) is applied to the full three-dimensional stage equation, with iteration matrix $I - hJ(Y)$ built from the full Jacobian J of Section 2.1, not from J_r directly. Section 3.3 shows that J_r is nonetheless the exact restriction of J to the invariant tangent space (Proposition 3), so the two matrices carry the same spectral information: any error in J_r would still be visible in the stiffness diagnostics of Section 4.2, even though the Newton solve itself uses J , not J_r .

3 Numerical Integration Methods

3.1 Explicit Euler Method

The explicit (forward) Euler method is the simplest one-step scheme for approximating the solution of an initial value problem [3, 4].

$$y'(t) = f(t, y(t)), \quad y(t_0) = y_0, \quad (19)$$

where, for the Robertson system, $y = (y_1, y_2, y_3)^\top$ and f is the right-hand side defined in Section 2.1. Given a step size $h > 0$ and a discrete time grid $t_n = t_0 + nh$, the method is obtained by replacing the derivative in (19) with a forward difference quotient, yielding the explicit update formula

$$y_{n+1} = y_n + h f(t_n, y_n), \quad n = 0, 1, 2, \dots \quad (20)$$

Because the right-hand side of (20) depends only on the already-known state y_n , no equation needs to be solved at each step; the new approximation y_{n+1} is computed directly. This non-implicit character is the defining feature of the method and the source of both its simplicity and its stability limitations, discussed in Section 4.

Local truncation error Expanding the exact solution $y(t_{n+1})$ in a Taylor series about t_n ,

$$y(t_{n+1}) = y(t_n) + h y'(t_n) + \frac{h^2}{2} y''(t_n) + \mathcal{O}(h^3), \quad (21)$$

and comparing with the update rule (20), the local truncation error (the error introduced in a single step, assuming the previous value y_n were exact) is

$$\tau_n = y(t_{n+1}) - y(t_n) - h f(t_n, y(t_n)) = \frac{h^2}{2} y''(t_n) + \mathcal{O}(h^3) = \mathcal{O}(h^2). \quad (22)$$

Since the local error is $\mathcal{O}(h^2)$ and accumulates over $\mathcal{O}(1/h)$ steps on a fixed interval $[t_0, T]$, the global error satisfies

$$\max_n \|y(t_n) - y_n\| = \mathcal{O}(h), \quad (23)$$

so explicit Euler is a first-order method. This low order motivates the introduction of the fourth-order Runge–Kutta method in Section 3.2 as a higher-accuracy alternative for the non-stiff regimes of the problem.

Application to the Robertson system For the Robertson system, (20) is applied componentwise to $y = (y_1, y_2, y_3)^\top$ with f as given in (5). Because the three reaction rate constants span roughly eight orders of magnitude, the fastest component of f dictates an extremely restrictive upper bound on h for (20) to remain numerically stable; this bound is derived quantitatively from the eigenvalues of the (reduced) Jacobian in Section 4.3. In the numerical experiments of Section 5, explicit Euler therefore serves primarily as a baseline that illustrates the practical consequences of stiffness, rather than as a competitive integrator for this problem.

3.2 Explicit Fourth-Order Runge–Kutta Method (RK4)

Method definition. The classical four-stage explicit Runge–Kutta method (RK4) is defined by the Butcher tableau

$$\begin{array}{c|ccc} 0 & & & \\ 1/2 & 1/2 & & \\ 1/2 & 0 & 1/2 & \\ 1 & 0 & 0 & 1 \\ \hline & 1/6 & 1/3 & 1/3 & 1/6 \end{array} \quad (24)$$

i.e. $c = (0, \frac{1}{2}, \frac{1}{2}, 1)$, $b = (\frac{1}{6}, \frac{1}{3}, \frac{1}{3}, \frac{1}{6})$, with the only nonzero sub-diagonal entries $a_{21} = \frac{1}{2}$, $a_{32} = \frac{1}{2}$, $a_{43} = 1$. For the autonomous system $y' = f(y)$ governing the Robertson problem, one step from y_n at t_n to y_{n+1} at $t_n + h$ reads

$$k_1 = f(y_n), \quad (25)$$

$$k_2 = f\left(y_n + \frac{h}{2} k_1\right), \quad (26)$$

$$k_3 = f\left(y_n + \frac{h}{2} k_2\right), \quad (27)$$

$$k_4 = f\left(y_n + h k_3\right), \quad (28)$$

$$y_{n+1} = y_n + \frac{h}{6} (k_1 + 2k_2 + 2k_3 + k_4). \quad (29)$$

Order conditions up to $p = 3$. We now derive, rather than quote, the conditions on (A, b, c) under which an explicit s -stage RK method reproduces the exact solution's Taylor expansion through $\mathcal{O}(h^3)$, and verify that Eq. (24) satisfies them.

For $y' = f(y)$, repeated use of the chain rule gives the exact Taylor expansion of the solution at y_n ,

$$y(t_n + h) = y_n + h f + \frac{h^2}{2} (Df \cdot f) + \frac{h^3}{6} [D^2 f(f, f) + Df \cdot Df \cdot f] + \mathcal{O}(h^4), \quad (30)$$

where $f = f(y_n)$, $Df = Df(y_n)$ is the Jacobian, and D^2f is the second (bilinear) derivative, all evaluated at y_n .

For a general explicit tableau, write $c_i = \sum_j a_{ij}$ (all stages start from $Y_i(0) = y_n$, so $k_i(0) = f$ for every i) and expand each stage as a power series in h ,

$$k_i(h) = f + h f_i^{(1)} + \frac{h^2}{2} f_i^{(2)} + O(h^3). \quad (31)$$

Since $Y_i(h) = y_n + h \sum_j a_{ij} k_j(h) = y_n + h c_i f + h^2 \sum_j a_{ij} f_j^{(1)} + O(h^3)$, Taylor-expanding $k_i(h) = f(Y_i(h))$ about y_n and matching powers of h gives

$$f_i^{(1)} = c_i (Df \cdot f), \quad f_i^{(2)} = 2 \left(\sum_j a_{ij} c_j \right) (Df \cdot Df \cdot f) + c_i^2 D^2 f(f, f). \quad (32)$$

Substituting into $y_{n+1} = y_n + h \sum_i b_i k_i(h)$ and collecting powers of h ,

$$\begin{aligned} y_{n+1} = y_n &+ h \left(\sum_i b_i \right) f + h^2 \left(\sum_i b_i c_i \right) (Df \cdot f) \\ &+ h^3 \left\{ \left(\sum_i b_i \sum_j a_{ij} c_j \right) (Df \cdot Df \cdot f) + \frac{1}{2} \left(\sum_i b_i c_i^2 \right) D^2 f(f, f) \right\} + O(h^4). \end{aligned} \quad (33)$$

Since f , $Df \cdot f$, $D^2 f(f, f)$, and $Df \cdot Df \cdot f$ are generically linearly independent elementary differentials, matching Eq. (33) to the exact expansion (30) term by term yields the order conditions for $p \leq 3$:

$$\sum_i b_i = 1, \quad (O1)$$

$$\sum_i b_i c_i = \frac{1}{2}, \quad (O2)$$

$$\sum_i b_i c_i^2 = \frac{1}{3}, \quad (O3a)$$

$$\sum_i b_i \sum_j a_{ij} c_j = \frac{1}{6}. \quad (O3b)$$

Condition (O1) matches the h^1 coefficient, (O2) the h^2 coefficient, and (O3a)–(O3b) together match the two independent h^3 terms — the coefficient of $D^2 f(f, f)$ and of $Df \cdot Df \cdot f$ respectively. Together these four scalar equations are necessary and sufficient for order $p = 3$; (O1)–(O2) alone give only order $p = 2$.

Proposition 4 (RK4 satisfies order $p \geq 3$). *The tableau (24) satisfies (O1)–(O3b).*

Proof. With $c = (0, \frac{1}{2}, \frac{1}{2}, 1)$ and $b = (\frac{1}{6}, \frac{1}{3}, \frac{1}{3}, \frac{1}{6})$:

$$\begin{aligned} \sum_i b_i &= \frac{1}{6} + \frac{1}{3} + \frac{1}{3} + \frac{1}{6} = 1, \\ \sum_i b_i c_i &= \frac{1}{3} \cdot \frac{1}{2} + \frac{1}{3} \cdot \frac{1}{2} + \frac{1}{6} \cdot 1 = \frac{1}{6} + \frac{1}{6} + \frac{1}{6} = \frac{1}{2}, \\ \sum_i b_i c_i^2 &= \frac{1}{3} \cdot \frac{1}{4} + \frac{1}{3} \cdot \frac{1}{4} + \frac{1}{6} \cdot 1 = \frac{1}{12} + \frac{1}{12} + \frac{2}{12} = \frac{1}{3}. \end{aligned}$$

For (O3b), only stages 3 and 4 contribute a nonzero $\sum_j a_{ij} c_j$ (stage 1 has no predecessors, and $a_{21} c_1 = \frac{1}{2} \cdot 0 = 0$ for stage 2): stage 3 gives $a_{32} c_2 = \frac{1}{2} \cdot \frac{1}{2} = \frac{1}{4}$, stage 4 gives $a_{43} c_3 = 1 \cdot \frac{1}{2} = \frac{1}{2}$. Hence

$$\sum_i b_i \sum_j a_{ij} c_j = b_3 \cdot \frac{1}{4} + b_4 \cdot \frac{1}{2} = \frac{1}{3} \cdot \frac{1}{4} + \frac{1}{6} \cdot \frac{1}{2} = \frac{1}{12} + \frac{1}{12} = \frac{1}{6}. \quad \square$$

Order conditions for $p = 4$. A fourth condition set, corresponding to the four rooted trees of order 4, is required for $p = 4$ and is not re-derived here for brevity:

$$\begin{aligned} \sum_i b_i c_i^3 &= \frac{1}{4}, & \sum_i b_i \left(\sum_j a_{ij} c_j \right)^2 & \text{(none; listed for completeness below)} \\ \sum_i b_i c_i \sum_j a_{ij} c_j &= \frac{1}{8}, & \sum_i b_i \sum_j a_{ij} c_j^2 &= \frac{1}{12}, \\ \sum_i b_i \sum_j a_{ij} \sum_k a_{jk} c_k &= \frac{1}{24}. \end{aligned} \tag{34}$$

(Listed as the standard eight order-4 conditions: the four in (O1)–(O3b) together with these four; see Hairer, Nørsett and Wanner [3], Thm. II.2.1.) Direct substitution of the RK4 tableau into all eight conditions confirms they hold exactly, which is the classical result that Eq. (29) is fourth order; we omit the remaining four substitutions here and cite the standard reference rather than reproduce them, as they add length without new insight beyond Proposition 4.

Application to the Robertson system RK4 is included as the high-accuracy *explicit* control group: it attains fourth-order accuracy at negligible per-step cost whenever the step size respects its stability restriction, unlike the implicit methods of Sections 3.3–3.4 — it has no mechanism for handling stiffness. Its bounded absolute stability region is derived in Section 4.1 and compared against the eigenvalue range of J_r from Section 2.3; Sections 5 and 6 show RK4 forced into the same prohibitively small step size as Explicit Euler once the fast manifold is active, despite its formally higher order. The observed-vs-theoretical convergence order for RK4 is verified empirically in Section 5.3.

3.3 Implicit Euler Method with Newton-Raphson Iteration

Stage equation and Newton iteration. One step of Implicit Euler, $y_{n+1} = y_n + hf(y_{n+1})$, requires solving the nonlinear stage equation [5]

$$g(Y) := Y - y_n - hf(Y) = 0 \tag{35}$$

for $Y \approx y_{n+1} \in \mathbb{R}^3$. We solve this with Newton–Raphson:

$$Y^{(k+1)} = Y^{(k)} - [Dg(Y^{(k)})]^{-1} g(Y^{(k)}), \quad Dg(Y) = I - hJ(Y), \tag{36}$$

where $J(Y) = Df(Y)$ is the full 3×3 Jacobian of the Robertson vector field (Section 2.1),

$$J(y) = \begin{pmatrix} -0.04 & 10^4 y_3 & 10^4 y_2 \\ 0.04 & -10^4 y_3 - 6 \times 10^7 y_2 & -10^4 y_2 \\ 0 & 6 \times 10^7 y_2 & 0 \end{pmatrix}. \tag{37}$$

The initial guess is the explicit-Euler predictor $Y^{(0)} = y_n + hf(y_n)$, and the iteration is judged converged when $\|g(Y^{(k)})\|_\infty \leq \text{tol}_N$ for a Newton tolerance tol_N set (and swept) in Section 5.5; Section 5.5 also tests two alternative initial guesses, the previous solution y_n and a deliberately perturbed predictor.

Relation to the reduced Jacobian of Section 2.3. $J(y)$ above and the 2×2 reduced Jacobian J_r derived in Section 2.3 are *not* the same matrix – they act on spaces of different dimension – but they are two representations of the same underlying linear map, related by the chain rule through the elimination $y_3 = 1 - y_1 - y_2$. Writing $\partial y_3 / \partial y_1 = \partial y_3 / \partial y_2 = -1$,

$$\left. \frac{\partial}{\partial y_i} \right|_{y_3 \text{ elim.}} = \frac{\partial}{\partial y_i} - \frac{\partial}{\partial y_3}, \quad i = 1, 2,$$

which is precisely the tangent-space projection established in Proposition 1 of Section 2.2 ($J(y)$ maps $\mathbb{R}^3$ into $T = \ker(\mathbf{1}^\top)$). Concretely, J_r is obtained from the first two rows of $J(y)$ by subtracting the third column from the first two – we verified this identity holds to floating-point precision by direct symbolic/numerical comparison. This is the sense in which the two Jacobians “are the same object”: the 2×2 form drives the eigenvalue/stiffness analysis of Section 2.3, while the full 3×3 form drives the Newton solve here; both are restrictions of one differential operator, so any qualitative statement about conditioning made from one is directly informative about the other.

Proposition 5 (exact conservation of the linear invariant under Newton). Let $\hat{Y}$ be any iterate produced by the Newton recursion above – converged or not – and define its residual $r := g(\hat{Y}) = \hat{Y} - y_n - hf(\hat{Y})$. Then

$$\mathbf{1}^\top \hat{Y} = \mathbf{1}^\top y_n + \mathbf{1}^\top r$$

holds as an exact algebraic identity, not an approximation.

Proof. By definition of r ,

$$\mathbf{1}^\top r = \mathbf{1}^\top \hat{Y} - \mathbf{1}^\top y_n - h \mathbf{1}^\top f(\hat{Y}).$$

Proposition 1 of Section 2.2 established $\mathbf{1}^\top f(y) \equiv 0$ for *every* $y \in \mathbb{R}^3$, not merely for y on the trajectory or satisfying the stage equation; in particular it holds at $y = \hat{Y}$ regardless of how close $\hat{Y}$ is to the true root. Hence the last term vanishes and rearranging gives the claim. $\square$

Corollary (tolerance-controlled conservation defect) If the Newton iteration is stopped once $\|r\|_\infty = \|g(\hat{Y})\|_\infty \leq \text{tol}_N$ (the stopping test of Section 5.5.1), and $\mathbf{1}^\top y_n = 1$ (true inductively from $\mathbf{1}^\top y_0 = 1$ and the Proposition applied at every prior step), then, using $\|r\|_2 \leq \sqrt{3} \|r\|_\infty$ and Cauchy–Schwarz,

$$|(\hat{y}_1 + \hat{y}_2 + \hat{y}_3) - 1| = |\mathbf{1}^\top r| \leq \sqrt{3} \|r\|_2 \leq 3 \|r\|_\infty \leq 3 \text{tol}_N.$$

The single-step conservation defect is therefore controlled linearly by the Newton tolerance – not by the time-step h or the method’s order – which is exactly what Section 5.4 verifies empirically, and exactly the quantity swept in Section 5.5’s Newton-tolerance study.

Numerical confirmation. We verified both statements directly on the Robertson vector field (not just the test equation): applying two Newton iterations from $Y^{(0)} = y_n = (1, 0, 0)$ with $h = 10^{-2}$ (deliberately *not* run to convergence, $\|g(\hat{Y})\| \approx 6.8 \times 10^{-2}$), the identity $\mathbf{1}^\top \hat{Y} - \mathbf{1}^\top y_n = \mathbf{1}^\top r$ held to within 3×10^{-17} (machine precision); running to full convergence ($\|g\| < 10^{-14}$) gave the same agreement, as expected since the identity does not depend on convergence.

Local convergence of Newton’s method. Once Newton’s method is sufficiently close to the exact nonlinear solution Y^* , its local error satisfies

$$\|Y^{(k+1)} - Y^*\| \leq C \|Y^{(k)} - Y^*\|^2$$

for some constant C , which is the standard quadratic convergence result for Newton iteration [6, 7].

One problem-specific caveat is worth flagging here rather than only in the results: since the spectrum of $J(y)$ contains eigenvalues of order 3×10^7 (Section 2.3), $Dg(Y) = I - hJ(Y)$ can become ill-conditioned for moderate h , shrinking the basin of quadratic convergence and potentially demanding more iterations – or a damped/globalized Newton step – near the fast transient at small t . We quantify this directly with an iteration-count and tolerance sweep in Section 5.5 rather than asserting it here.

3.4 Trapezoidal Method (Crank–Nicolson) as a Second Implicit Scheme

To complement the first-order Implicit Euler scheme and establish a multi-order comparison framework, we introduce the classical Trapezoidal method as a second implicit integrator [5]. The method discretizes the autonomous system $\frac{dy}{dt} = f(y)$ by averaging the vector fields at the current and next time steps:

$$y_{n+1} = y_n + \frac{h}{2} [f(y_n) + f(y_{n+1})], \quad (38)$$

where $h > 0$ denotes the integration step size.

Non-linear Residual and Newton–Raphson Formulation Following the iterative framework established in Section 3.3, Eq. (38) defines an implicit algebraic system for the unknown state vector y_{n+1} . We recast this relation into finding the root of the non-linear residual vector function $G : \mathbb{R}^d \rightarrow \mathbb{R}^d$:

$$G(y_{n+1}) \equiv y_{n+1} - y_n - \frac{h}{2} f(y_n) - \frac{h}{2} f(y_{n+1}) = 0. \quad (39)$$

Solving $G(y_{n+1}) = 0$ via the multivariate Newton–Raphson method yields the linear update step at iteration k :

$$W(y_{n+1}^{(k)}) \Delta y^{(k)} = -G(y_{n+1}^{(k)}), \quad y_{n+1}^{(k+1)} = y_{n+1}^{(k)} + \Delta y^{(k)}, \quad (40)$$

where the iteration matrix $W(y) \in \mathbb{R}^{d \times d}$ corresponds to the Fréchet derivative $DG(y)$:

$$W(y) \equiv DG(y) = \frac{\partial G}{\partial y}(y) = I - \frac{h}{2} J_f(y). \quad (41)$$

Here, $I \in \mathbb{R}^{d \times d}$ is the identity matrix, and $J_f(y) \equiv \frac{\partial f}{\partial y}(y)$ represents either the full 3×3 Jacobian or the reduced 2×2 Jacobian matrix $J_r(\tilde{y})$ derived in Section 2.3.

Second-Order Accuracy By symmetric Taylor expansion of the stage evaluations around $t_n + \frac{h}{2}$, the local truncation error (LTE) of the Trapezoidal method satisfies:

$$\tau_{n+1} = y(t_{n+1}) - y_{n+1} = \mathcal{O}(h^3). \quad (42)$$

Accumulating LTE over $N = T/h$ integration steps yields a *second-order global discretization error*:

$$E_{\text{global}}(h) = \max_{0 \leq n \leq N} \|y(t_n) - y_n\| = \mathcal{O}(h^2). \quad (43)$$

This offers a higher order of accuracy ($\mathcal{O}(h^2)$) compared to the first-order ($\mathcal{O}(h)$) convergence of the Implicit Euler scheme.

Role in Multi-Order Cross-Consistency Verification The inclusion of the second-order Trapezoidal scheme serves an essential structural role in our overall numerical validation strategy. Together with the first-order Implicit Euler method, it forms the foundation for evaluating numerical self-consistency across methods of different orders—representing the third component of our threefold validation protocol detailed in Section 5.3.

Remark on Asymptotic Behavior: While the Trapezoidal scheme achieves second-order accuracy, its asymptotic damping characteristics under severe stiffness differ fundamentally from Implicit Euler. The rigorous stability analysis regarding its lack of L -stability (despite being A -stable) is detailed in Section 4.1.

3.5 Adaptive Step-Size Control

To improve the efficiency of Explicit Euler on the stiff Robertson problem, a simple step-doubling controller is implemented together with the other integrators, following standard adaptive time-stepping ideas [8, 9].

Error estimate. From $\mathbf{y}_n$ with step h , one full step and two half-steps are computed:

$$\begin{aligned}\mathbf{y}_{n+1}^{(1)} &= \mathbf{y}_n + h \mathbf{f}(\mathbf{y}_n), \\ \mathbf{y}_{n+1/2} &= \mathbf{y}_n + \frac{h}{2} \mathbf{f}(\mathbf{y}_n), \quad \mathbf{y}_{n+1}^{(2)} = \mathbf{y}_{n+1/2} + \frac{h}{2} \mathbf{f}(\mathbf{y}_{n+1/2}),\end{aligned}$$

and the local error is estimated by

$$e_n = \|\mathbf{y}_{n+1}^{(2)} - \mathbf{y}_{n+1}^{(1)}\|_\infty.$$

The local error of a single Euler step is Ch^2 and that of two half-steps is about $Ch^2/2$, so the difference is, to leading order, the local error of $\mathbf{y}_{n+1}^{(2)}$; the Richardson factor $1/(2^p - 1) = 1$ for $p = 1$ therefore needs no rescaling. The more accurate value is kept, $\mathbf{y}_{n+1} = \mathbf{y}_{n+1}^{(2)}$.

Acceptance and step-size update. A step is accepted when $e_n \leq \text{tol}$ (or when $h \leq 1.01 h_{\min}$); otherwise it is rejected and repeated with a smaller step. The next step is

$$h_{\text{new}} = h s \left(\frac{\text{tol}}{e_n} \right)^{1/(p+1)}, \quad p = 1, \quad s = 0.9,$$

with the factor h_{new}/h limited to $[0.5, 2]$ after an accepted step (to 2 if $e_n = 0$) and to $[0.1, 0.5]$ after a rejected one, and h_{new} clipped to $[h_{\min}, h_{\max}]$. The parameters are $h_0 = 10^{-4}$, $\text{tol} = 10^{-6}$, $h_{\min} = 10^{-10}$, $h_{\max} = 10^{-2}$. The tolerance is absolute and applied to the maximum norm over the three components; since $y_2 = O(10^{-5})$, it does not control y_2 in a relative sense (see the failure at $\text{tol} = 10^{-4}$ in Section 5.5.2).

Demonstration run ($\text{tol} = 10^{-6}$). Over $0 \leq t \leq 40$ the controller accepts 25 398 steps and rejects 12 686 trial steps (33% of all trials). The accepted step sizes range from 1.0×10^{-4} to 8.3×10^{-3} , with mean 1.57×10^{-3} and median 1.73×10^{-3} , so the step is far from uniform. In the initial transient it grows quickly: after 11 accepted steps ($t \approx 7 \times 10^{-3}$) it exceeds 3×10^{-3} . It does *not*, however, keep growing where the solution is smooth: the median step is 1.9×10^{-3} for $1 \leq t < 10$ and only 1.4×10^{-3} for $t \geq 10$, decreasing to about 9×10^{-4} near $t = 40$, although y_2 then varies slowly (Figure 1a).

The final state is

$$\mathbf{y}(40) \approx (0.715813151, 9.14185 \times 10^{-6}, 0.284177707),$$

against the reference $\mathbf{y}_{\text{ref}}(40) \approx (0.715827069, 9.1855 \times 10^{-6}, 0.284163746)$ of Section 5.1. The final-state error is $\|\mathbf{y}(40) - \mathbf{y}_{\text{ref}}(40)\|_2 \approx 1.97 \times 10^{-5}$; the maximum over the whole trajectory is $E_\infty = 1.97 \times 10^{-5}$ and is attained at the end. The invariant defect is 1.33×10^{-14} and no negative component occurs.




Figure 1: Adaptive Explicit Euler, $\text{tol} = 10^{-6}$, every accepted step shown. (a) y_2 against the reference. (b) Accepted step sizes against the local stability bounds $2/|\lambda_{\max}(t)|$ (Explicit Euler) and $4/|\lambda_{\max}(t)|$ (two half-steps), with $|\lambda_{\max}|$ evaluated on the reference trajectory (Section 4.2). (c) Local error estimate e_n of every accepted step in units of the tolerance.

Local versus global error. Panel (c) of Figure 1 shows that the controller does what it is designed to do locally: the estimate of every accepted step lies below the tolerance (maximum 0.97 tol), and for 90% of the steps between 0.19 and 0.81 of it, i.e. within a factor of about five. The global error is nevertheless about 20 times the tolerance. The tolerance bounds a per-step estimate, whereas the global error accumulates over 25 398 steps; the sum of the local estimates is 1.3×10^{-2} (at most $N \text{ tol} = 2.5 \times 10^{-2}$) and is much larger than the actual global error, which indicates that the stiff dynamics damp most of the local errors, but that the global error is not controlled by tol .

Why the step does not grow. Panel (b) explains the behaviour of h . For $t \gtrsim 0.2$ the accepted steps alternate between two levels: a lower one below $4/|\lambda_{\max}|$, whose two half-steps of size $h/2$ lie inside the Explicit-Euler stability interval, and an upper one above it, with local error estimates of about 0.2 and 0.8 of the tolerance (panel c). For $t > 0.1$ the accepted steps lie between 1.5 and 3.0 times the Explicit-Euler limit $2/|\lambda_{\max}(t)|$ (5th–95th percentile). The controller therefore rides the linear stability boundary of the explicit method and stays

there only through rejections, which are about one third of all trial steps. Adaptive step-size control is thus an efficiency improvement for an explicit method, not a way around stiffness; its dependence on the tolerance is quantified in Section 5.5.2, and Section 5.6 compares it with an adaptive implicit solver.

4 Linear Stability Analysis

This section collects the local (linear) stability tools and applies them to the Robertson system. Section 4.1 derives the absolute-stability regions of the four integrators for the scalar test equation; Section 4.2 computes the Jacobian eigenvalues, the stiffness ratio and a non-normality diagnostic along our own reference trajectory; Section 4.3 turns these into a step-size prediction for the explicit methods and compares it with the experiments of Section 5.4. Throughout, “local” means that the Jacobian is frozen at a state of the reference trajectory.

4.1 Absolute Stability Regions of the Four Methods

To characterise each integrator independently of the nonlinear system, we apply it to the scalar test equation $y' = \lambda y$, $\lambda \in \mathbb{C}$, and set $z = h\lambda$. Standard absolute-stability theory is used here [10, 5]. One step of any of the four methods of Section 3 reads $y_{n+1} = R(z) y_n$ with a method-specific growth factor $R(z)$, and the region of absolute stability is $\mathcal{S} = \{z \in \mathbb{C} : |R(z)| \leq 1\}$.

Explicit Euler. $y_{n+1} = (1 + z)y_n$, so $R(z) = 1 + z$. The region $|1 + z| \leq 1$ is the closed unit disk centred at $z = -1$; on the real axis it gives $z \in [-2, 0]$.

RK4. For an explicit Runge–Kutta method of order p applied to the test equation, $R(z)$ is the degree- p Taylor polynomial of e^z ; for classical RK4,

$$R(z) = 1 + z + \frac{z^2}{2} + \frac{z^3}{6} + \frac{z^4}{24}.$$

On the negative real axis R is positive and not monotone: it decreases from $R(0) = 1$ to a minimum $R \approx 0.270$ at $z \approx -1.60$ and then increases again, returning to $R = 1$ at $z \approx -2.7853$ (root of $R(x) = 1$ found with Brent’s method). The real-axis stability interval is therefore $[-2.7853, 0]$, about 1.4 times that of Explicit Euler but still bounded. Away from the real axis the region has the bulging shape of Figure 2 (top right).

Implicit Euler. $y_{n+1} = y_n + z y_{n+1}$, hence $R(z) = 1/(1 - z)$. *A-stability.* For $z = x + iy$ with $x \leq 0$, $|1 - z|^2 = (1 - x)^2 + y^2 = 1 - 2x + x^2 + y^2 \geq 1$, so $|R(z)| \leq 1$ on the closed left half-plane. In fact $\mathcal{S} = \{z : |z - 1| \geq 1\}$ is strictly larger: only the open unit disk centred at $z = 1$ is excluded (Figure 2, bottom left). *L-stability.* $R(z) \rightarrow 0$ as $z \rightarrow -\infty$ (numerically $R(-10^8) \approx 1.0 \times 10^{-8}$), so the fastest decaying modes are damped strongly: Implicit Euler is A-stable and L-stable.

Trapezoidal / Crank–Nicolson. $y_{n+1} = y_n + \frac{h}{2}\lambda(y_n + y_{n+1})$, hence $R(z) = \frac{1 + z/2}{1 - z/2}$. *A-stability, exactly.* With $w = z/2 = u + iv$, $|1 + w| \leq |1 - w| \iff (1 + u)^2 \leq (1 - u)^2 \iff u \leq 0$. The stability region is exactly the closed left half-plane, and on the imaginary axis $|R(iy)| = 1$ (numerically $||R(iy)| - 1| \leq 4 \times 10^{-16}$ for $10^{-3} \leq y \leq 10^3$). *Not L-stable.* $R(z) \rightarrow -1$ as $z \rightarrow -\infty$ (numerically $R(-10^8) \approx -1.0000$). Very fast decaying modes are therefore reflected with amplitude close to one and alternating sign instead of being damped. This is a concrete, testable prediction; Section 5.6 confirms that Trapezoidal/CN loses non-negativity at large steps ($h = 5 \times 10^{-3}$ and 10^{-2}) although it is A-stable.




Figure 2: Absolute-stability regions $|R(z)| \leq 1$ (shaded) of the four integrators on the common domain $\text{Re } z \in [-6, 3]$, $\text{Im } z \in [-4, 4]$.

Caveat. These regions describe the scalar, linear, autonomous test equation only; they say nothing directly about the nonlinear Robertson system. A nonlinear theory exists in the form of Dahlquist’s B-stability [5], which guarantees $|y_{n+1} - z_{n+1}| \leq |y_n - z_n|$ under the one-sided Lipschitz condition $\langle f(u) - f(v), u - v \rangle \leq 0$, but the condition is not established for the Robertson vector field, which contains indefinite quadratic terms. A-/L-stability are therefore used as necessary local diagnostics, to be corroborated by the eigenvalue analysis of Section 4.2 and the step-size sweep of Section 5.4.

4.2 Eigenvalue Analysis of the Robertson System

Setup. The Jacobian is evaluated along our own reference trajectory (Radau, $\text{rtol} = 10^{-12}$, $\text{atol} = 10^{-14}$, Section 5.1) through its dense output, on the 201 output times ($t = 0$ and 200 log-spaced points in $[10^{-8}, 40]$) and, for smooth curves, on a 1000-point logarithmic grid on the same interval.

Structural zero and initial degeneracy. Because $\mathbf{1}^T \mathbf{f} \equiv 0$, the full 3×3 Jacobian satisfies $\mathbf{1}^T J(\mathbf{y}) = \mathbf{0}^T$ for every state: it has the eigenvalue 0 for all t (left eigenvector $\mathbf{1}$), and a stiffness ratio that includes it is infinite or undefined. Eliminating $y_3 = 1 - y_1 - y_2$ gives the reduced Jacobian $J_r = [J_{ij} - J_{i3}]_{i,j=1,2}$ of Section 2.2. The plane $\mathbf{1}^T \mathbf{v} = 0$ is invariant under J , so the two eigenvalues of J_r are exactly the two nonzero eigenvalues of J ; we confirmed this numerically (largest relative difference 6×10^{-15} over all sampled states). We therefore define

$$S(t) = \frac{\max\{|\lambda_{\text{slow}}|, |\lambda_{\text{fast}}|\}}{\min\{|\lambda_{\text{slow}}|, |\lambda_{\text{fast}}|\}}$$

from the eigenvalues of J_r . If the full Jacobian is used instead, an eigenvalue is classified as zero when $|\lambda| < 10^{-9}$: along the trajectory the structural zero has modulus at most 1×10^{-15} , while the smallest nonzero eigenvalue is 0.0214 (at $t = 40$), so the threshold is unambiguous.

The initial state $(1, 0, 0)$ is a *different*, genuine degeneracy: with $y_2 = y_3 = 0$ the fast reaction terms vanish, $J(0)$ has a double zero eigenvalue, and $J_r(0) = \begin{pmatrix} -0.04 & 0 \\ 0.04 & 0 \end{pmatrix}$ has a zero eigenvalue in addition to -0.04 . The ratio $S(0)$ is therefore undefined and all curves start at the first positive output time $t = 10^{-8}$, where S is finite ($S \approx 1.67$).

Results. Figure 3 shows $|\lambda_{\text{slow}}|$, $|\lambda_{\text{fast}}|$ and $S(t)$. Two features deserve comment. (i) For small t the growing eigenvalue is $\lambda \approx -2k_2y_2 \approx -2k_2k_1t$, while the other stays near $-k_1 = -0.04$. The two moduli coincide at $t_c = 1/(2k_2) \approx 1.67 \times 10^{-8}$; for $t < t_c$ the ordering by modulus is reversed, so $S = \max/\min$ touches 1 at t_c (a property of the definition, not of the dynamics). (ii) For $t \gtrsim 10^{-3}$, $|\lambda_{\text{fast}}|$ saturates near 2.2×10^3 and later grows to 3.393×10^3 at $t = 40$ as y_3 increases, whereas $|\lambda_{\text{slow}}| \approx 0.4$ decays to 0.0214; hence S has a plateau near 5.4×10^3 and rises to 1.58×10^5 at $t = 40$.




Figure 3: Top: magnitudes of the two nonzero eigenvalues of $J_r(t)$ along the reference trajectory. Bottom: stiffness ratio $S(t)$, starting at the first positive output time; markers are the orientation times of Table 1. The dotted line marks $t_c = 1/(2k_2)$.

Table 1: Stiffness ratio $S(t)$ from the reduced Jacobian along our own reference, evaluated at the exact orientation times, compared with the Problem-Pack check values (which are quoted to two or three digits and used only as a self-consistency target).

t	our $S(t)$	Problem-Pack value	relative difference
1×10^{-4}	2.99×10^3	3.0×10^3	0.23%
1×10^{-2}	5.42×10^3	5.4×10^3	0.42%
40	1.58×10^5	1.58×10^5	0.25%

Table 1 compares $S(t)$ at the three orientation times (evaluated exactly, not at the nearest grid point) with the Problem-Pack check values; the differences are below 0.5%, i.e. within the two to three digits in which those values are quoted. At $t = 40$, $|\lambda|_{\max} = 3.393 \times 10^3$ and

$$J_r(40) \approx \begin{pmatrix} -0.1319 & 2841.5 \\ 0.1319 & -3392.7 \end{pmatrix}.$$

Non-normality. J_r is not symmetric, so eigenvalues alone can mislead. For a diagonalisable $J = V\Lambda V^{-1}$, one Explicit-Euler step matrix satisfies

$$\|(I + hJ)^n\|_2 \leq \kappa(V) \max_j |1 + h\lambda_j|^n, \quad \kappa(V) = \|V\|_2 \|V^{-1}\|_2,$$

so $\kappa(V)$ bounds how much the eigenvalue-based factor can understate transient amplification ($\kappa = 1$ for a normal matrix). Figure 4 shows $\kappa(V(t))$: on the 201 output times its median is 1.014, only 2% of the points exceed 10, and $\kappa = 2.14$ at $t = 40$. The peak (54.3 at $t = 1.74 \times 10^{-8}$ on the output grid) is a resolution artefact: it sits at the eigenvalue crossing t_c , where the two eigenvalues coincide and J_r is nearly a Jordan block, so the peak grows as the grid is refined (229 on the 1000-point grid). Where it occurs $|\lambda| \approx 0.04$, so $h|\lambda| \ll 2$ for every step size used and non-normality cannot influence stability. At late times, where $h|\lambda_{\text{fast}}| \approx 2$ decides stability, $\kappa \lesssim 2.2$, so the frozen-Jacobian estimate of the next section is reliable to within that factor.




Figure 4: Condition number $\kappa(V(t))$ of the eigenvector matrix of $J_r(t)$. Dots: the 201 output times; line: 1000-point grid. The dotted line marks t_c .

4.3 Step-Size Constraints: Theory versus Experiment

Prediction. For a real negative eigenvalue the Explicit-Euler condition $|1 + h\lambda| \leq 1$ gives $h \leq 2/|\lambda|$. Freezing the Jacobian at the state $\mathbf{y}_{\text{ref}}(t)$, the local restriction is $h \lesssim 2/|\lambda_{\text{fast}}(t)|$, and a single step size valid on all of $[0, 40]$ must satisfy

$$h < \frac{2}{\max_{t \in [0, 40]} |\lambda_{\text{fast}}(t)|} = \frac{2}{|\lambda|_{\max}} = \frac{2}{3.393 \times 10^3} \approx 5.895 \times 10^{-4}, \quad (44)$$

the maximum being attained at $t = 40$. The same argument with the RK4 real-axis limit gives $h < 2.7853/|\lambda|_{\max} \approx 8.21 \times 10^{-4}$. Implicit Euler and Trapezoidal/CN are A-stable, so linear stability imposes no restriction on h ; accuracy and admissibility still have to be checked (Section 5.4).




Figure 5: Local $h\lambda$ diagnostic on the real axis: $h|\lambda_{\text{fast}}(t)|$ along the reference trajectory for four step sizes, against the real-axis limits of Explicit Euler (2) and RK4 (2.7853). A curve above a limit means that the frozen-Jacobian condition is violated at that time.

Figure 5 shows where the condition fails. For $h = 5 \times 10^{-4}$ it holds throughout; for $h = 6 \times 10^{-4}$, 7×10^{-4} and 8×10^{-4} it is first violated at $t \approx 37.0$, 18.2 and 8.6. The RK4 step $h = 8 \times 10^{-4}$ used in the convergence study reaches only 97% of its limit at $t = 40$ and stays below it.

Table 2: Explicit Euler: time at which the local limit $h|\lambda_{\text{fast}}(t)| > 2$ is first exceeded along the reference trajectory, compared with the outcome of the step-size sweep of Section 5.4.

h	local limit exceeded from $t =$	observed outcome	divergence time
4.00×10^{-4}	never	admissible	—
5.00×10^{-4}	never	admissible	—
5.50×10^{-4}	never	admissible	—
5.90×10^{-4}	39.9	admissible	—
5.95×10^{-4}	38.4	admissible	—
6.00×10^{-4}	37.0	reaches T , negative values	—
7.00×10^{-4}	18.2	diverged	27.1
8.00×10^{-4}	8.6	diverged	17.7

Experiment. The step-size sweep is reported in Section 5.4 (Table 5) and is not repeated here: all steps up to 5.95×10^{-4} are admissible, 6.0×10^{-4} reaches T with negative concentrations, and 7×10^{-4} , 8×10^{-4} diverge; a bisection places the boundary at $h_{\text{emp}} \approx 5.962 \times 10^{-4}$, i.e. 1.1% above h_{crit} . Table 2 sets the local prediction beside these outcomes. In both diverging runs the blow-up occurs about nine time units after the local limit is first exceeded (8.9 and 9.1), which is plausible because a mildly unstable mode needs time to grow from round-off and truncation errors. For $h = 6 \times 10^{-4}$ the limit is exceeded only from $t \approx 37.0$, i.e. too late for a comparable delay to end in divergence before $T = 40$, in agreement with the observed finite but negative solution. The same reasoning explains why $h = 5.9 \times 10^{-4}$ and 5.95×10^{-4} , slightly above h_{crit} , are still admissible. This delay was measured on two runs only, so it supports the picture but does not prove it.

Interpretation and limitations. The agreement should be read as empirical support for a local estimate, not as a nonlinear stability theorem: the Jacobian varies along the trajectory, and the scalar argument describes only the linearised amplification. Non-normality is mild where the constraint is decided ($\kappa \lesssim 2.2$ at late times). Finally, linear stability does not by itself guarantee positivity or conservation: at $h = 6 \times 10^{-4}$ the invariant defect is still at round-off level while the solution is already negative, so the three properties have to be tested separately (Section 5.4).

5 Numerical Experiments and Validation

This section asks whether the numerical results for the Robertson problem can be trusted, and what they reveal about stiffness. Three questions structure the experiments: (i) are the implementations correct, i.e. do they reproduce their formal orders against a verified reference? (ii) which step-size restriction is active for each method, accuracy or stability, and is the solution still physically admissible? (iii) what does each method cost at comparable accuracy? Stability, conservation, non-negativity and accuracy are treated as *separate* properties: a trajectory can be finite and conserve $y_1 + y_2 + y_3 = 1$ while being inaccurate or negative. The numerical comparison follows the standard distinction between methods intended for nonstiff and stiff ODE systems [11, 12]. In particular, stiff-oriented implicit algorithms are designed to avoid the severe stability restrictions that can dominate explicit integration.

The four fixed-step integrators and the step-doubling adaptive Euler method are implemented from scratch using NumPy arrays [13], with the time-stepping loops compiled using Numba [14]; Newton’s method uses the analytic Jacobian and a dense 3×3 linear solve. All numerical figures were generated using Matplotlib [15]. SciPy’s Radau and BDF [16, 17] are used only as the independent oracle for the reference and, in Section 5.6, as a library baseline. All errors use

$$E_\infty = \max_n \|\mathbf{y}_n - \mathbf{y}_{\text{ref}}(t_n)\|_2, \quad (45)$$

where t_n runs over *every* time point stored by the method and $\mathbf{y}_{\text{ref}}$ is evaluated by dense output. The accuracy of that dense output is checked in Section 5.1.

5.1 Reference Solution Construction

The reference is SciPy’s Radau method [16, 5] with $\text{rtol} = 10^{-12}$, $\text{atol} = 10^{-14}$. Two independent checks support it: the run is repeated with tolerances 100 times looser, and a BDF run with the same tight tolerances is compared with it. Both comparisons are made at the 201 prescribed output times ($t = 0$ and 200 log-spaced points in $[10^{-8}, 40]$) and on a fine grid of 60,000 points, which tests the dense output used in Eq. (45).

Table 3: Validation checks for the reference solution (maximum Euclidean discrepancy).

Check	Compared at	Discrepancy
Radau, tolerances tightened by 100	201 output times	6.30×10^{-12}
Radau vs. BDF (same tolerances)	201 output times	1.94×10^{-11}
Radau, tolerances tightened by 100	60,000 fine-grid points (dense output)	1.03×10^{-11}
Radau vs. BDF	60,000 fine-grid points (dense output)	1.95×10^{-11}
Invariant defect $\max y_1 + y_2 + y_3 - 1 $	201 output times	1.33×10^{-15}

All discrepancies are about 2×10^{-11} or smaller (Table 3), so only nine decimals are retained:

$$\mathbf{y}_{\text{ref}}(40) \approx (0.715827069, 9.1855 \times 10^{-6}, 0.284163746).$$

The smallest error reported below ($E_\infty = 1.2 \times 10^{-10}$, RK4) is only about six times the Radau–BDF discrepancy, i.e. at the edge of what this reference can certify; no smaller error is claimed.




Figure 6: Tight Radau reference trajectory (values below $\text{atol} = 10^{-14}$ are not drawn). y_2 stays several orders of magnitude below y_1 and y_3 .

5.2 Sanity Checks

Before interpreting convergence, stability, and efficiency results, we verify several problem-specific sanity checks. These checks reuse the reference-validation, physical-admissibility, and Newton-tolerance experiments reported in this section.

The reference solution is independently supported by Radau tolerance refinement, a Radau–BDF comparison, and the invariant defect reported in Table 3. In particular,

$$\max_t |y_1(t) + y_2(t) + y_3(t) - 1| = 1.33 \times 10^{-15},$$

and

$$y_{\text{ref}}(40) \approx (0.715827069, 9.1855 \times 10^{-6}, 0.284163746).$$

Physical admissibility is checked separately from accuracy. The fixed-step runs in Table 4 remain non-negative, whereas the Explicit Euler boundary study in Table 5 shows that a step size of 6.0×10^{-4} already produces negative concentrations even though the invariant defect remains near round-off. This confirms that conservation alone is not sufficient to establish a physically valid solution. The expected transient behaviour is also observed: y_2 remains small, peaks early, and subsequently decays to a small positive value at $t = 40$.

For the implicit methods, Table 6 verifies that the Newton tolerance is sufficiently strict. Loose tolerances produce non-physical solutions or Newton failure, while sufficiently tight tolerances preserve the invariant and prevent nonlinear-solver error from dominating the reported time-discretisation error.

5.3 Convergence Order Verification

Validation protocol. Because no exact trajectory is available, correctness rests on three independent checks. (i) *Independent oracle*: the Radau reference, verified by a tolerance refinement and by BDF (Section 5.1). (ii) *Known properties*: the linear invariant $y_1 + y_2 + y_3 = 1$ and non-negativity, tracked separately from accuracy (Section 5.4). (iii) *Cross-order self-consistency*: methods of orders 1, 2 and 4 must converge to the same limit with their formal orders. At their finest step sizes the RK4 and Trapezoidal/CN runs differ from the reference by 1.2×10^{-10} and

4.9×10^{-8} , so by the triangle inequality they agree with each other to within 5×10^{-8} ; the two Euler methods, of first order, reach 6.7×10^{-7} (explicit) and 1.3×10^{-5} (implicit) at their finest steps, consistent with their orders. The slopes are extracted below.

For a method of formal order p , $E(h) \approx Ch^p$, so the observed pairwise order between successive step sizes is

$$p_{\text{obs}} = \frac{\log(E_i/E_{i+1})}{\log(h_i/h_{i+1})}. \quad (46)$$

Each method is run on four successive halvings inside its stable range. Explicit Euler starts at $h = 5 \times 10^{-4}$, just below its stability limit (Section 5.4); RK4 starts at $h = 8 \times 10^{-4}$, which is already 97% of its own real-axis limit $h \lesssim 8.21 \times 10^{-4}$ (the RK4 stability interval is $[-2.785, 0]$). The implicit runs solve the nonlinear system to $\tau_N = 10^{-13}$ (Section 5.5.1).

Table 4: Four-method convergence study. All 16 runs reached $T = 40$ with finite values (stable in this sense); “Non-neg.” states whether $\min_{n,j} y_{j,n} \geq -10^{-12}$; p_{fit} is the least-squares slope over the three finest step sizes.

Method	h	E_∞	Pairwise order	Non-neg.	p_{fit}
Explicit Euler	5.00×10^{-4}	6.831×10^{-6}	—	yes	1.060
	2.50×10^{-4}	2.893×10^{-6}	1.240	yes	
	1.25×10^{-4}	1.365×10^{-6}	1.084	yes	
	6.25×10^{-5}	6.655×10^{-7}	1.036	yes	
RK4	8.00×10^{-4}	1.172×10^{-6}	—	yes	4.247
	4.00×10^{-4}	4.377×10^{-8}	4.743	yes	
	2.00×10^{-4}	2.179×10^{-9}	4.329	yes	
	1.00×10^{-4}	1.215×10^{-10}	4.165	yes	
Implicit Euler	1.00×10^{-2}	1.039×10^{-4}	—	yes	1.000
	5.00×10^{-3}	5.200×10^{-5}	0.999	yes	
	2.50×10^{-3}	2.601×10^{-5}	1.000	yes	
	1.25×10^{-3}	1.301×10^{-5}	1.000	yes	
Trapezoidal / CN	1.00×10^{-3}	1.470×10^{-6}	—	yes	2.004
	5.00×10^{-4}	7.825×10^{-7}	0.910	yes	
	2.50×10^{-4}	1.949×10^{-7}	2.005	yes	
	1.25×10^{-4}	4.864×10^{-8}	2.003	yes	

Using one rule for all methods—the least-squares slope over the three finest step sizes—gives $p_{\text{EE}} \approx 1.06$, $p_{\text{RK4}} \approx 4.25$, $p_{\text{IE}} \approx 1.000$ and $p_{\text{CN}} \approx 2.004$ (Table 4, Figure 7). The coarsest pairs are pre-asymptotic: CN gives 0.910, and RK4’s first pair (4.74) lies close to its stability limit; the pairwise orders then decrease towards the formal values (Explicit Euler $1.24 \rightarrow 1.04$, RK4 $4.74 \rightarrow 4.17$). The table also records that every run is finite, reaches $T = 40$ and is non-negative, so the differences seen here are purely differences of accuracy.




Figure 7: Convergence orders with respect to E_∞ . The dashed, dash-dotted and dotted lines are slope guides of order 1, 2 and 4.

5.4 Explicit Euler: Stability Boundary and Physical Admissibility

The eigenvalue analysis of Section 4.2 gives $|\lambda|_{\max} = 3.393 \times 10^3$, attained at the final time $t = 40$, and therefore the frozen-Jacobian estimate of Eq. (44)

$$h_{\text{crit}} = \frac{2}{|\lambda|_{\max}} \approx 5.895 \times 10^{-4}.$$

This is a local estimate, not a nonlinear stability theorem, so it is tested by a step-size sweep in which three diagnostics are kept apart: non-negativity, full-state error, and the invariant defect $\max_n |y_{1,n} + y_{2,n} + y_{3,n} - 1|$.

Table 5: Explicit Euler near the predicted limit h_{crit} . Admissible: reaches T with $\min y_j \geq -10^{-12}$; negative: reaches T but violates non-negativity; diverged: blow-up before T (time at which $\max_j |y_j|$ first exceeds 2). The invariant defect is evaluated only over steps with $\max_j |y_j| \leq 2$, because beyond that it measures round-off of huge numbers.

h	Status	Minimum y_j	Invariant defect	Final-time error
4.00×10^{-4}	admissible	0	1.07×10^{-14}	1.98×10^{-6}
5.00×10^{-4}	admissible	0	1.24×10^{-14}	2.47×10^{-6}
5.50×10^{-4}	admissible	0	2.62×10^{-14}	2.72×10^{-6}
5.90×10^{-4}	admissible	0	1.07×10^{-14}	2.91×10^{-6}
5.95×10^{-4}	admissible	0	1.20×10^{-14}	2.94×10^{-6}
6.00×10^{-4}	negative values	-4.84×10^{-6}	6.66×10^{-15}	4.48×10^{-3}
7.00×10^{-4}	diverged ($t \approx 27.1$)	—	6.66×10^{-15}	—
8.00×10^{-4}	diverged ($t \approx 17.7$)	—	4.33×10^{-15}	—

Table 5 and Figure 8 show that all steps up to $h = 5.95 \times 10^{-4}$ are admissible. At $h = 6.0 \times 10^{-4}$ the run still reaches $T = 40$, but concentrations become negative ($\min y_j = -4.8 \times 10^{-6}$) and the final error grows by three orders of magnitude to 4.5×10^{-3} ; at $h = 7 \times 10^{-4}$ and 8×10^{-4} it diverges at $t \approx 27$ and $t \approx 18$. A bisection between 5.95×10^{-4} and 6.0×10^{-4} places the admissibility boundary at $h_{\text{emp}} \approx 5.962 \times 10^{-4}$, i.e. 1.1% above h_{crit} (the bisection assumes a monotone transition inside this interval). The small excess is consistent with the local nature of the estimate: since $|\lambda_{\max}|$ is largest at the end, the local limit $2/|\lambda_{\max}(t)|$ is exceeded by h_{emp} only for $t \gtrsim 38.1$, so the weakly amplified fast mode has little time to grow (Table 2 in Section 4.3 lists these times for all step sizes). For this trajectory the frozen-Jacobian value is thus an accurate and slightly conservative predictor; this is evidence, not a proof.

The three diagnostics fail independently. The invariant defect stays at round-off level [18] ($\sim 10^{-14}$) even at $h = 6.0 \times 10^{-4}$, where the solution is already negative and inaccurate: every explicit Runge–Kutta step preserves $\mathbf{1}^T \mathbf{y}$ because $\mathbf{1}^T \mathbf{f} \equiv 0$, so conservation carries no information about admissibility or accuracy. (For the diverged runs the defect is evaluated only while $\max_j |y_j| \leq 2$; beyond that it merely measures round-off of huge numbers.)




Figure 8: Explicit Euler near h_{crit} . (a) final-time error; (b) minimum concentration over the run. The step size is scaled by h_{crit} ; the shaded region contains the diverged runs $h = 7 \times 10^{-4}$, 8×10^{-4} (Table 5).

5.5 Solver Diagnostics

5.5.1 Newton-Tolerance Sensitivity and Initial Guess

Implicit Euler solves a nonlinear system in every step, so its error contains time-discretisation error and algebraic error. To separate them, $h = 10^{-2}$ is fixed and the stopping tolerance τ_N of the *absolute* residual test $\|F\|_\infty \leq \tau_N$ is varied. Because $y_2 = O(10^{-5})$ (peak 3.6×10^{-5}), τ_N has to be judged against y_2 , not against the $O(1)$ components.

Table 6: Newton-tolerance sensitivity of Implicit Euler at fixed $h = 10^{-2}$ (4000 steps). The stopping test is $\|F\|_\infty \leq \tau_N$ on the absolute residual; “Newton updates” counts linear solves summed over all steps.

τ_N	Status	$y_2(40)$	$ \Delta y_2 $	Final L^2 error	Newton updates	Max residual	Invariant defect
1×10^{-2}	non-physical	-3.591×10^{-6}	1.28×10^{-5}	3.63×10^1	7860	1.0×10^{-2}	1.3×10^{-13}
1×10^{-4}	failed ($t = 3.71$)	—	—	—	—	—	8.9×10^{-16}
1×10^{-6}	completed	9.191×10^{-6}	5.79×10^{-9}	5.44×10^{-5}	2095	1.0×10^{-6}	1.8×10^{-15}
1×10^{-8}	completed	9.187×10^{-6}	1.16×10^{-9}	4.41×10^{-5}	3427	1.0×10^{-8}	2.9×10^{-15}
1×10^{-10}	completed	9.187×10^{-6}	1.36×10^{-9}	4.94×10^{-5}	4012	9.4×10^{-12}	1.6×10^{-15}
1×10^{-12}	completed	9.187×10^{-6}	1.36×10^{-9}	4.94×10^{-5}	4014	1.1×10^{-13}	2.4×10^{-15}

Table 6 shows three regimes. For $\tau_N = 10^{-2}$ and 10^{-4} , which exceed the peak of y_2 , the solve cannot resolve y_2 : in both runs a component turns negative from $t \approx 0.02$ on ($\min y_j = -28.8$ and -8.0), and the $\tau_N = 10^{-4}$ run finally fails at $t \approx 3.71$, when Newton no longer converges within 50 iterations from the unphysical state. A loose tolerance therefore does not merely lose accuracy; it destroys admissibility, and the failure appears long after the damage was done. For $\tau_N \leq 10^{-6}$ the solution is physical, but the error is converged only for $\tau_N \leq 10^{-10}$: the final error is 5.44×10^{-5} (10^{-6}) and 4.41×10^{-5} (10^{-8}) against the converged 4.94×10^{-5} , and $|\Delta y_2|$ settles at 1.36×10^{-9} . These deviations are consistent with per-step residuals accumulating over the 4000 steps (worst case $N\tau_N = 4 \times 10^{-5}$ for $\tau_N = 10^{-8}$), which suggests the rule $\tau_N \lesssim E_{target}/N_{steps}$. For the fixed-step study, with up to 6.4×10^5 steps, $\tau_N = 10^{-13}$ was chosen

so that algebraic error cannot contaminate the order estimates; we did not test whether a looser value would suffice there.

The invariant defect does *not* diagnose an unconverged Newton solve: it is $\lesssim 10^{-13}$ for every τ_N in Table 6. Since $\mathbf{1}^T J \equiv 0$, a single Newton update with the exact Jacobian restores $\sum_j y_j$ to its value at the previous step, so residual size and defect are decoupled here; the residual, y_2 and the error must be monitored instead.

The initial guess matters as well. At $h = 10^{-2}$ and $\tau_N = 10^{-10}$, the explicit-Euler predictor and the previous solution both converge to the same physical solution ($E_\infty = 1.04 \times 10^{-4}$), whereas a deliberately perturbed predictor ($\Delta y_1 = +0.1$, $\Delta y_2 = -0.1$) passes the same residual test but converges to a non-physical root of the implicit equation ($E_\infty = 7.06$, 7652 negative entries). A small residual therefore certifies neither correctness nor admissibility.

5.5.2 Adaptive Step-Size Control

The adaptive method is step-doubling Explicit Euler: the error estimate is $\|y_{\text{full}} - y_{\text{two halves}}\|_\infty \leq \text{tol}$, the two-half-step value is kept, and the step is rescaled by $0.9\sqrt{\text{tol}/\text{err}}$ (limited to $[0.5, 2]$ after acceptance and $[0.1, 0.5]$ after rejection). Being explicit, it is validated by tolerance refinement, not by a fixed-step convergence line.

Table 7: Adaptive Explicit Euler (step doubling, $h_0 = 10^{-4}$, $h_{\max} = 10^{-2}$). The tolerance bounds the *local* error estimate per step; E_∞ is the accumulated *global* error.

Tol.	Status	Accepted	Rejected	E_∞	E_∞/tol	Mean accepted h
1×10^{-4}	diverged ($t \approx 0.06$)	416	40	—	—	—
1×10^{-5}	completed	22248	9794	1.02×10^{-3}	102	1.80×10^{-3}
3×10^{-6}	completed	25279	12544	1.32×10^{-4}	44	1.58×10^{-3}
1×10^{-6}	completed	25398	12686	1.97×10^{-5}	20	1.57×10^{-3}
3×10^{-7}	completed	25407	12692	1.20×10^{-5}	40	1.57×10^{-3}
1×10^{-7}	completed	23202	9194	1.35×10^{-5}	135	1.72×10^{-3}
3×10^{-8}	completed	25221	12145	1.24×10^{-5}	413	1.59×10^{-3}
1×10^{-8}	completed	23794	8854	9.67×10^{-6}	967	1.68×10^{-3}

The step size does adapt (Figure 1, Section 3.5): at $\text{tol} = 10^{-6}$ it grows from $h_0 = 10^{-4}$ to as much as 8.3×10^{-3} during the initial transient and then settles between about 9×10^{-4} and 3×10^{-3} , following the slowly growing $|\lambda_{\max}(t)|$. What it does *not* do is respect the linear stability limit: for $t > 0.1$ the accepted steps lie between 1.5 and 3.0 times $2/|\lambda_{\max}(t)|$ (5th–95th percentile), straddling the limit $4/|\lambda_{\max}|$ for the two half-steps. The method survives only because about 33% of all trial steps are rejected—consistent with the fast mode being amplified until the step-doubling estimate exceeds the tolerance.

The global error does not track the tolerance (Table 7): E_∞/tol is already 20 at 10^{-6} and reaches 967 at 10^{-8} , while E_∞ saturates near 10^{-5} , about the accuracy of fixed-step Explicit Euler at its stability limit. Two effects act together: the tolerance controls a *local* estimate whereas E_∞ accumulates over more than 2×10^4 steps, and below $\text{tol} \approx 10^{-6}$ stability, not accuracy, sets the step (the step count stays near 2.5×10^4). A tolerance that is too loose fails outright: $\text{tol} = 10^{-4}$ exceeds the peak of y_2 , cannot see its growth, and diverges at $t \approx 0.06$. The requirement “error close to the tolerance” is therefore not met by an explicit method on this stiff problem, which is the point of the experiment; Section 5.6 shows what an adaptive *implicit* method achieves.

5.6 Efficiency Comparison

Efficiency is compared at comparable accuracy, with the runtime measured as the median of seven warmed calls of the compiled kernels (plotting excluded). Times depend on machine and

implementation; the step and work counts are the reproducible measures. Figure 9 extends the step-size range of the implicit methods beyond the convergence study, and hollow markers flag runs with negative concentrations.




Figure 9: Work-precision diagram. Hollow markers: $\min y_j < 0$. The star is SciPy’s adaptive Radau at $\text{rtol} = 10^{-6}$ (library baseline).

Table 8: Cost at comparable accuracy (median of warmed runs, this machine). Times are implementation- and machine-dependent; the SciPy row is not JIT-compiled and its time is not like-for-like with the compiled kernels, so compare its work counts.

Method	h	E_∞	Min. y_j	Time (ms)	Work
<i>(a) Target $E_\infty \leq 10^{-4}$: fastest tested run per method</i>					
Explicit Euler	5.00×10^{-4}	6.83×10^{-6}	≥ 0	32.1	80000 steps; 80000 RHS evaluations
Implicit Euler	5.00×10^{-3}	5.20×10^{-5}	≥ 0	7.7	8000 steps; 8014 Newton updates
RK4	8.00×10^{-4}	1.17×10^{-6}	≥ 0	25.8	50000 steps; 200000 RHS evaluations
Trapezoidal / CN	2.00×10^{-3}	2.81×10^{-6}	≥ 0	21.1	20000 steps; 20015 Newton updates
Trapezoidal / CN (non-physical)	1.00×10^{-2}	9.41×10^{-5}	-2.9×10^{-5}	3.7	4000 steps; 4102 Newton updates
<i>(b) Matched accuracy $E_\infty \approx 7 \times 10^{-6}$</i>					
Explicit Euler	5.00×10^{-4}	6.83×10^{-6}	≥ 0	32.1	80000 steps; 80000 RHS evaluations
Implicit Euler	6.50×10^{-4}	6.76×10^{-6}	≥ 0	62.9	61539 steps; 61558 Newton updates
<i>(c) Library reference: adaptive implicit solver</i>					
SciPy Radau, $\text{rtol} = 10^{-6}$	–	6.05×10^{-8}	≥ 0	12.5	66 steps; 545 RHS evaluations; 14 Jacobians

Target $E_\infty \leq 10^{-4}$. Explicit Euler cannot enlarge its step (Section 5.4), so its best point is 15 times more accurate than required: the accuracy is wasted, which is the practical signature of stiffness. Implicit Euler meets the target with $h = 5 \times 10^{-3}$ in the fewest steps and least time among the admissible runs. RK4 is confined to $h \lesssim 8.21 \times 10^{-4}$ and cannot exploit its order beyond that limit. Trapezoidal/CN reaches the target fastest with $h = 10^{-2}$, but with negative

concentrations ($\min y_j = -2.9 \times 10^{-5}$): the rule is A-stable but not L-stable, so the fast mode is not damped at large h . In our runs $h = 5 \times 10^{-3}$ is also inadmissible, whereas $h \leq 2 \times 10^{-3}$ is admissible, and there CN costs about as much as Explicit Euler or RK4.

Matched accuracy $E_\infty \approx 7 \times 10^{-6}$. Here Implicit Euler needs $h = 6.5 \times 10^{-4}$ (61539 steps) and takes about 2.0 times as long as Explicit Euler in this environment. For the first-order pair the implicit method therefore pays off only when the accuracy demand is loose enough for its step to exceed h_{crit} by a large factor: in Figure 9 the Implicit-Euler curve drops below the cost of the best Explicit-Euler run (about 15 ms) only for $E_\infty \gtrsim 2 \times 10^{-5}$; for tighter accuracy the cheaper explicit step wins.

Adaptive implicit baseline. SciPy’s Radau with $\text{rtol} = 10^{-6}$ reaches $E_\infty \approx 6.0 \times 10^{-8}$ in 66 steps (545 right-hand-side evaluations)—about 400 times fewer steps and an error about 300 times smaller than adaptive Euler at the same tolerance. Its wall-clock time is not directly comparable with the compiled kernels, but its work counts are.

Summary of findings. (1) The Radau reference is reproduced by a tolerance refinement and by BDF to about 2×10^{-11} , including its dense output. (2) All four implementations recover their formal orders ($\approx 1.06, 4.25, 1.000, 2.004$). (3) Stiffness appears as a *stability* restriction, not an accuracy restriction: Explicit Euler is admissible only up to $h \approx 5.962 \times 10^{-4}$ (the linear limit of RK4 is 8.21×10^{-4}), while Implicit Euler reaches $E_\infty \approx 10^{-4}$ with $h = 10^{-2}$, roughly seventeen times larger; adaptivity does not remove the restriction for an explicit method. (4) Conservation, non-negativity, Newton residual and error are independent diagnostics, and each can hide the failure of the others. (5) The implicit advantage depends on the accuracy demanded (in our environment it appears only for $E_\infty \gtrsim 2 \times 10^{-5}$), and for CN it requires checking non-negativity.

6 Discussion

6.1 Physical Interpretation via Quasi-Steady-State Approximation (QSSA)

The numerical results of Sections 4 and 5 raise two questions that a purely numerical study cannot answer: why does the intermediate species B stay so small, and why does the problem behave as a stiff one even though B is almost absent? The reference solution (Figure 6) shows y_2 climbing from zero to about 3.6×10^{-5} within a few 10^{-3} time units, remaining near this level for more than two decades in time and then decaying slowly to $y_2(40) = 9.1855 \times 10^{-6}$, although B is consumed with rate constants up to $k_2 = 3 \times 10^7$. This section explains the magnitude and the time structure of y_2 by a quasi-steady-state argument [19, 20], makes it rigorous with a singular-perturbation (Tikhonov) framework, and connects it to the eigenvalue structure of Section 4.2 and to the numerical findings of Section 5.

Order-of-magnitude balance. The equation for the intermediate species is $\dot{y}_2 = k_1 y_1 - k_3 y_2 y_3 - k_2 y_2^2$. At early times $y_1 \approx 1$ and $y_3 \approx 0$, so $\dot{y}_2 \approx k_1 - k_2 y_2^2$: B is produced slowly (rate $k_1 = 0.04$) and consumed by a quadratic term with an enormous coefficient. Setting $\dot{y}_2 = 0$ gives the balance $k_1 y_1 \approx k_2 y_2^2$, i.e.

$$y_2^* \approx \sqrt{\frac{k_1 y_1}{k_2}} = \sqrt{\frac{0.04}{3 \times 10^7}} \approx 3.65 \times 10^{-5}. \quad (47)$$

The reference has $\max_t y_2 = 3.6487 \times 10^{-5}$, attained at $t = 4.56 \times 10^{-3}$, only 0.08% below Eq. (47). B stays small because the quadratic sink reaches the production rate at a tiny concentration, and $y_2^* \propto k_2^{-1/2}$ decreases as the sink becomes stronger. Two points refine the picture “a spike followed by a rapid relaxation”. First, there is no overshoot: $\max_t y_2$ lies slightly *below* y_2^* , so the rise is a monotone approach to the quasi-steady value. Second, what

follows the rise is not a fast relaxation but a slow decay on the time scale of the slow variables, along which y_2 follows its quasi-steady value.

One balance, two regimes. As y_3 grows, the sink $k_3 y_2 y_3$ takes over from $k_2 y_2^2$ and the balance becomes $k_1 y_1 \approx k_3 y_2 y_3$, i.e. $y_2^* \approx k_1 y_1 / (k_3 y_3)$. Both are limits of the single algebraic balance $k_2 y_2^2 + k_3 y_3 y_2 = k_1 y_1$, whose positive root is

$$y_2^{\text{QSSA}}(y_1, y_3) = \frac{-k_3 y_3 + \sqrt{k_3^2 y_3^2 + 4k_1 k_2 y_1}}{2k_2}. \quad (48)$$

Table 9 and Figure 10b compare the three predictions with the reference, evaluated with the reference values of y_1, y_3 . The early balance is accurate at $t = 10^{-2}$ (error at most 0.16% on $5 \times 10^{-3} < t < 10^{-2}$) but is 17% off at $t = 1$ and a factor 3.4 too large at $t = 40$; the late balance is the opposite and is 9.7% off at $t = 40$ (the value 1.01×10^{-5} against 9.1855×10^{-6}). Eq. (48) is accurate throughout: its relative error is below 9.2×10^{-5} for all $t \geq 5 \times 10^{-3}$ and 3.2×10^{-6} at $t = 40$.

Table 9: Quasi-steady predictions of y_2 divided by the reference value $y_2(t)$, evaluated with the reference y_1, y_3 . Early balance: $\sqrt{k_1 y_1 / k_2}$; late balance: $k_1 y_1 / (k_3 y_3)$; quadratic QSSA: positive root of $k_2 y_2^2 + k_3 y_3 y_2 = k_1 y_1$.

t	reference y_2	early balance	late balance	quadratic QSSA
1×10^{-2}	3.645×10^{-5}	1.0016	302	0.99991
1	3.075×10^{-5}	1.17	3.75	0.99993
40	9.186×10^{-6}	3.36	1.10	1.00000

Initial layer. For $t \lesssim 10^{-3}$ one has $y_1 \approx 1$ and $y_3 \approx 0$, and the equation $\dot{y}_2 = k_1 - k_2 y_2^2$ with $y_2(0) = 0$ can be solved exactly:

$$y_2(t) = \sqrt{\frac{k_1}{k_2}} \tanh(\sqrt{k_1 k_2} t), \quad t_f = \frac{1}{\sqrt{k_1 k_2}} = 9.13 \times 10^{-4}. \quad (49)$$

For small t this reduces to $y_2 \approx k_1 t$, the linear growth already seen in Section 4.2. Eq. (49) predicts that y_2 reaches 99% of y_2^* at $t = 2.42 \times 10^{-3}$; the reference gives 2.43×10^{-3} . Its relative error is at most 3.7×10^{-5} for $t \leq 10^{-3}$ (Figure 10a). The “spike” is thus a boundary layer of thickness $\sim t_f$, three to four orders of magnitude shorter than the slow time scale $1/k_1 = 25$. Adding the layer correction to Eq. (48), $y_2 \approx y_2^{\text{QSSA}} - y_2^*(0) [1 - \tanh(t/t_f)]$, brings the relative error below 9.2×10^{-5} on the whole interval (Figure 10c).




Figure 10: Quasi-steady-state analysis of y_2 . (a) Initial layer on $0 \leq t \leq 8 \times 10^{-3}$: reference, Eq. (49) and the early balance y_2^* . (b) Reference against the early balance, the late balance and the quadratic QSSA, Eq. (48). (c) Relative error of the quadratic QSSA and of the QSSA with initial-layer correction; the dash-dotted line marks t_f . All predictions use the reference y_1, y_3 .

Singular-perturbation framework (Tikhonov). A naive choice of the small parameter $\varepsilon = 1/k_2$ gives $\varepsilon \dot{y}_2 = \varepsilon(k_1 y_1 - k_3 y_2 y_3) - y_2^2$, whose critical manifold $y_2 = 0$ is degenerate ($\partial g / \partial y_2 = -2y_2 = 0$ there), so Tikhonov's theorem does not apply; this reflects that the quasi-steady value is $O(\sqrt{\varepsilon})$, not $O(\varepsilon)$. The remedy is to rescale $y_2 = \delta u$ with

$$\delta = k_2^{-1/2} = 1.83 \times 10^{-4}, \quad \kappa_3 = k_3 \delta = 1.83 = O(1),$$

where $k_3 \approx 10^4$ is regarded as being of order $k_2^{1/2}$. Since $k_2 \delta^2 = 1$, the system takes the standard slow-fast form

$$\dot{y}_1 = -k_1 y_1 + \kappa_3 u y_3, \quad \dot{y}_3 = u^2, \quad \delta \dot{u} = k_1 y_1 - \kappa_3 u y_3 - u^2 =: g(y_1, y_3, u), \quad (50)$$

with the invariant $y_1 + y_3 + \delta u = 1$. Tikhonov's theorem [21, 22] requires three conditions, all satisfied here:

- (i) *Critical manifold.* $g = 0$ has the isolated positive root $u = \varphi(y_1, y_3) = \frac{1}{2}[-\kappa_3 y_3 + \sqrt{\kappa_3^2 y_3^2 + 4k_1 y_1}]$, which is exactly Eq. (48) divided by δ .
- (ii) *Normal attraction.* $\partial g / \partial u = -2u - \kappa_3 y_3$, which on the manifold equals $-\sqrt{\kappa_3^2 y_3^2 + 4k_1 y_1} \leq -2\sqrt{k_1 y_1} < 0$ uniformly along the trajectory ($y_1 \geq 0.71$).
- (iii) *Initial value in the domain of attraction.* With $\tau = t/\delta$, the boundary-layer equation is $du/d\tau = k_1 - u^2$ (for $y_1 = 1, y_3 = 0$); every $u > -\sqrt{k_1}$ is attracted to $\sqrt{k_1}$, and $u(0) = 0$ qualifies. Its solution is Eq. (49).

The theorem then states that, for every fixed $t_1 > 0$, $(y_1, y_3) = (\bar{y}_1, \bar{y}_3) + O(\delta)$ and $u = \varphi(\bar{y}_1, \bar{y}_3) + O(\delta)$ uniformly on $[t_1, T]$, where $(\bar{y}_1, \bar{y}_3)$ solves the reduced system obtained by inserting $u = \varphi$ into the first two equations of (50). The deviation from the manifold decays like $\exp(-\mu t/\delta)$ with $\mu = |\partial g / \partial u|$, i.e. within a time $O(\delta)$ (t_f in Eq. (49)).

We checked the statement quantitatively. The reduced model (solved with Radau, without any use of the full solution) differs from the full solution by at most 1.3×10^{-5} in y_1 and 3.6×10^{-5} in y_3 for $t \geq 10^{-2}$; both are below $\delta = 1.83 \times 10^{-4}$, as predicted. The offset in y_3 is not random: the reduced model starts on the manifold, so it produces y_3 too early, by $k_2 \int_0^\infty [y_2^{*2} - y_2^2] dt = k_1 t_f = 3.65 \times 10^{-5}$, which is what is observed. This is the boundary-layer correction of the outer solution.

Link to the eigenvalue structure of Section 4.2. Condition (ii) also gives the fast eigenvalue: on the manifold $\lambda_{\text{fast}} = (\partial g / \partial u) / \delta = -(2k_2 y_2 + k_3 y_3)$, i.e.

$$\lambda_{\text{fast}} = -\sqrt{k_3^2 y_3^2 + 4k_1 k_2 y_1}. \quad (51)$$

This explains the three stages seen in Figure 3. (i) *Sharp early separation:* before y_2 reaches its quasi-steady value, $\lambda_{\text{fast}} \approx -2k_2 y_2 \approx -2k_1 k_2 t$ grows linearly in t , so the stiffness ratio rises from 1 to about 5×10^3 by $t \approx 10^{-3}$, i.e. in the initial layer. (ii) *Plateau:* once y_2 is slaved, $\lambda_{\text{fast}} \approx -2\sqrt{k_1 k_2} = -2.191 \times 10^3$ ($y_3 \approx 0$), against the computed 2.190×10^3 at $t = 10^{-2}$. The fast rate saturates near 2×10^3 , and not at $k_2 = 3 \times 10^7$, because k_2 multiplies the tiny y_2 , not 1; this is the reason for the remark that $|\lambda|_{\text{max}} = O(10^3)$. (iii) *Late milder growth:* as y_3 increases, Eq. (51) rises only to 3.3928×10^3 at $t = 40$ (computed 3.3928×10^3). Over more than four decades in time the fast eigenvalue therefore changes by a factor of about 1.5; the stiffness ratio nevertheless grows again to 1.6×10^5 because the slow eigenvalue decays from 0.4 to 0.02. Eq. (51) also gives the explicit-Euler limit in closed form: $h_{\text{crit}} = 2/\sqrt{k_3^2 y_3(T)^2 + 4k_1 k_2 y_1(T)} = 5.895 \times 10^{-4}$, identical to Eq. (44).

Consequences for the numerical results.

- *Stiffness is a stability, not an accuracy, restriction* (Sections 5.4 and 5.6). After about $3t_f \approx 3 \times 10^{-3}$, i.e. less than 0.01% of the interval, the fast mode carries no signal: y_2 is a slaved variable. Accuracy needs to resolve only the slow dynamics, but Explicit Euler must still keep $|1 + h\lambda_{\text{fast}}| < 1$ for the (numerically excited) fast mode, hence $h < 5.895 \times 10^{-4}$; this is why its best run is about 15 times more accurate than required. The adaptive controller of Section 3.5 cannot escape this limit either (Section 5.5.2).
- *L-stability decides what happens to departures from the manifold* (Section 4.1). For Implicit Euler $R(z) \rightarrow 0$: a deviation from the slow manifold is removed within one step, and the method effectively integrates the reduced model, which is why $h = 10^{-2}$ is admissible.

For Trapezoidal/CN $R(z) \rightarrow -1$: an unresolved deviation is not damped but reflected with alternating sign. At $h = 10^{-2}$ and 5×10^{-3} the layer, of duration t_f , is not resolved, and the observed negative concentrations ($\min y_j = -2.9 \times 10^{-5}$ and -4.2×10^{-5} , Section 5.6) have the size of $y_2^* = 3.65 \times 10^{-5}$ itself, which is what an undamped, unresolved layer perturbation of amplitude $\sim y_2^*$ would produce. We regard this as consistent with, but not a proof of, the mechanism.

- *Tolerances must resolve y_2^** (Sections 5.5.1 and 5.5.2). For an implicit step the Newton solve determines y_2 through the algebraic balance $g = 0$. An absolute residual tolerance of 10^{-2} or 10^{-4} , larger than y_2^* , cannot determine y_2 , which is why those runs turned negative; the adaptive run with tolerance 10^{-4} , also larger than the peak of y_2 , diverged.

Limits of the picture. The quasi-steady description is valid only after the layer: at $t = 2 \times 10^{-3}$ the QSSA alone is still 2.5% off. Tikhonov's theorem guarantees only $O(\delta)$ accuracy; the observed relative accuracy of about 10^{-4} after the correction is much better than this bound, which we do not claim to have explained. Finally, Figure 10 and Table 9 use the reference values of y_1, y_3 and are therefore diagnostics of the balance; the independent test of the reduced model is the comparison above.

6.2 Why Explicit Methods Fail and Implicit Methods Succeed

The evidence of Sections 4–5 forms one causal chain.

Cause. After the initial layer of duration $t_f = 9.13 \times 10^{-4}$ the fast mode is slaved to the slow dynamics (Section 6.1), but it stays in the Jacobian: $|\lambda_{\text{fast}}| = \sqrt{k_3^2 y_3^2 + 4k_1 k_2 y_1}$ rises from 2.2×10^3 to 3.393×10^3 , while the slow eigenvalue is only 0.02–0.4 and the stiffness ratio reaches 1.6×10^5 (Section 4.2).

Prediction. The stability regions of Explicit Euler and RK4 are bounded, with real-axis intervals $[-2, 0]$ and $[-2.7853, 0]$ (Section 4.1). With $|\lambda|_{\text{max}}$ attained at $t = 40$ this gives $h < 5.895 \times 10^{-4}$ and $h < 8.21 \times 10^{-4}$ (Section 4.3). Implicit Euler and Trapezoidal/CN are A-stable and impose no linear restriction; only Implicit Euler is also L-stable ($R \rightarrow 0$ instead of $R \rightarrow -1$).

Test. The sweep of Section 5.4 confirms the prediction: Explicit Euler is admissible up to $h_{\text{emp}} = 5.962 \times 10^{-4}$ (1.1% above h_{crit}), becomes negative at 6×10^{-4} and diverges at 7×10^{-4} and 8×10^{-4} , about nine time units after the local limit is first exceeded (Table 2). RK4 at $h = 8 \times 10^{-4}$ (97% of its limit) is stable and accurate but cannot use a larger step, so its fourth order is of no help. The adaptive controller rides the same limit, at the price of about one third rejected steps (Sections 3.5 and 5.5.2).

Consequence. Explicit methods do not fail for lack of accuracy: their best run is 15 times more accurate than required, because the step is dictated by a mode that is already extinct. Implicit Euler reaches $E_\infty \approx 10^{-4}$ with $h = 10^{-2}$, about seventeen times h_{crit} . The success is conditional, however: the method is first order, and in our environment it is cheaper than the best Explicit-Euler run only for $E_\infty \gtrsim 2 \times 10^{-5}$; at matched accuracy 7×10^{-6} Explicit Euler is about 2.0 times faster (Section 5.6).

The Trapezoidal prediction. Section 4.1 predicted that Trapezoidal/CN, although A-stable, reflects very fast modes ($R \rightarrow -1$) instead of damping them. With $\lambda_{\text{fast}} \approx -2.19 \times 10^3$, the growth factor at $h = 10^{-3}, 2 \times 10^{-3}, 5 \times 10^{-3}, 10^{-2}$ is $R_{\text{CN}} \approx -0.05, -0.37, -0.69, -0.83$,

whereas $R_{\text{IE}} \approx 0.31, 0.19, 0.08, 0.04$: a perturbation left by an unresolved layer is nearly preserved with alternating sign by CN and removed by Implicit Euler. This is what Section 5.6 observes: CN has negative concentrations at $h = 5 \times 10^{-3}$ and 10^{-2} ($\min y_j = -4.2 \times 10^{-5}$ and -2.9×10^{-5} , the size of $y_2^* = 3.65 \times 10^{-5}$) but not at $h \leq 2 \times 10^{-3}$. We did not see the effect as a visible oscillation in an error curve: its amplitude is comparable to the error of Implicit Euler, so E_∞ hides it (CN at $h = 10^{-2}$ has $E_\infty = 9.4 \times 10^{-5}$, below 1.04×10^{-4} for Implicit Euler); it shows only in the sign of y_2 .

6.3 Three-Fold Validation Strategy Recap

Our confidence in the reported numbers rests on three independent chains of evidence.

(i) Independent oracle (Section 5.1). The Radau reference agrees with a run at 100 times looser tolerances to 6.30×10^{-12} and with BDF to 1.94×10^{-11} at the 201 output times, and to 1.95×10^{-11} on a 60,000-point grid, which validates the dense output used for every E_∞ (Table 3).

(ii) Known properties (Sections 5.4, 4.1 and 6.1). Analytic facts are compared with computation: the invariant $y_1 + y_2 + y_3 = 1$ (Section 2.2) is preserved to round-off ($\sim 10^{-14}$) by the admissible runs we checked, while non-negativity and accuracy are checked separately; the stability regions and the RK4 limit 2.7853 are exact results; the initial-layer solution of Section 6.1 matches the reference to 3.7×10^{-5} ; and the stiffness ratio agrees with the Problem-Pack check values to their quoted digits (Table 1).

(iii) Cross-order self-consistency (Section 5.3). Methods of orders 1, 2 and 4 converge to the same reference with fitted slopes 1.06, 1.000 (first order), 2.004 (second) and 4.25 (fourth); the finest RK4 and CN runs, with errors 1.2×10^{-10} and 4.9×10^{-8} , agree with each other to within 5×10^{-8} .

The prediction–observation loop of Section 4.3 (h_{crit} against h_{emp}) connects these chains to the stability claims. This is the chain of evidence on which we base the final numbers: $\mathbf{y}(40) \approx (0.715827069, 9.1855 \times 10^{-6}, 0.284163746)$ with no error claimed below 2×10^{-11} . The chain verifies the numerics, not the model: the oracle and all methods solve the same ODE, whose coefficients are fixed by the Problem Pack.

6.4 Limitations

- **Frozen-Jacobian analysis (Sections 4.1 and 4.3).** $h_{\text{crit}} = 2/|\lambda|_{\text{max}}$ is a local estimate along the *reference* trajectory, using real parts only; a numerical trajectory that has already left it (e.g. after the first negative value) sees other Jacobians. h_{emp} comes from a bisection that assumes a monotone transition, and the nine-time-unit delay between exceeding the local limit and divergence was measured on two runs only.
- **Non-normality (Section 4.2).** We computed only $\kappa(V)$ of the 2×2 reduced Jacobian, whose peak is a resolution artefact; the bound $\kappa(V) \max_j |1 + h\lambda_j|^n$ holds for a frozen matrix. We did not compute pseudospectra or the product of the step matrices along the trajectory, so transient growth is bounded only heuristically.
- **Newton iteration (Section 3.3).** The solver is a full Newton method with the analytic Jacobian, an Explicit-Euler predictor and an absolute residual test, without damping or positivity safeguards. Its convergence region was not analysed, only probed: a perturbed guess converges to a non-physical root, and a loose tolerance fails at $t \approx 3.71$ (Section 5.5.1). The value $\tau_N = 10^{-13}$ used for the order study was chosen conservatively and not tested against looser ones.

- **Adaptive control (Section 3.5).** The controller is step-doubling Explicit Euler with an absolute max-norm tolerance and no stability safeguard; the global error is about 20 times the tolerance. An adaptive *implicit* method was not implemented; SciPy’s Radau serves only as a library baseline whose wall time is not comparable with the compiled kernels.
- **Scope of the empirical conclusions.** All results concern one initial value problem, one parameter set and $T = 40$. Timings, the crossover near $E_\infty \approx 2 \times 10^{-5}$ and the factor 2.0 come from one machine and one implementation; the reference certifies errors only down to $\sim 2 \times 10^{-11}$; and Tikhonov’s theorem guarantees only $O(\delta)$, whereas the QSSA is observed to be far more accurate, which we do not explain (Section 6.1).

7 Conclusions

We asked whether numerical solutions of the Robertson problem on $0 \leq t \leq 40$ can be trusted, and what they reveal about stiffness. Four self-implemented integrators (Explicit Euler, RK4, Implicit Euler and Trapezoidal/CN with Newton’s method) and a step-doubling adaptive Euler method were compared with a tight Radau reference, whose discrepancy with BDF and with a looser run is about 2×10^{-11} . Our main findings are as follows.

1. **The results are trustworthy.** The implementations recover their formal orders (fitted slopes 1.06, 4.25, 1.000, 2.004), and three independent chains of evidence support the numbers (Section 6.3): $\mathbf{y}(40) \approx (0.715827069, 9.1855 \times 10^{-6}, 0.284163746)$.
2. **Stiffness is a stability restriction, not an accuracy restriction.** The fast mode dies within $t_f \approx 9.13 \times 10^{-4}$ and y_2 is then slaved to a quasi-steady value (Section 6.1), yet it keeps $|\lambda|_{\max}$ at $O(10^3)$, not 3×10^7 . The frozen-Jacobian estimate $h_{\text{crit}} = 5.895 \times 10^{-4}$ predicts the Explicit-Euler boundary $h_{\text{emp}} = 5.962 \times 10^{-4}$ to 1.1%; RK4 is limited to 8.21×10^{-4} , and step-size adaptivity does not remove the limit.
3. **Implicit methods succeed, with a caveat.** L-stable Implicit Euler reaches $E_\infty \approx 10^{-4}$ with $h = 10^{-2}$, about seventeen times h_{crit} . Trapezoidal/CN is A-stable but not L-stable, and loses non-negativity for $h \geq 5 \times 10^{-3}$, as predicted.
4. **Diagnostics must be kept apart.** Conservation stays at round-off even when the solution is negative or inaccurate; a small Newton residual can belong to a non-physical root; tolerances must resolve $y_2 \sim 10^{-5}$.
5. **Efficiency depends on the accuracy demanded.** In our environment the implicit advantage appears only for $E_\infty \gtrsim 2 \times 10^{-5}$; below that Explicit Euler is cheaper. An adaptive implicit solver (SciPy Radau) needs 66 steps for $E_\infty \approx 6.0 \times 10^{-8}$, against 25 398 for adaptive Euler at a larger error.

Practical recommendation. For Robertson-type problems, use an L-stable implicit or a stiff adaptive solver, choose the Newton tolerance well below the smallest relevant component, and monitor non-negativity and error in addition to conservation. Explicit methods are appropriate only when the accuracy target is so tight that the step is limited by accuracy anyway.

Limitations and outlook. The conclusions are empirical and restricted to one problem, one machine and a frozen-Jacobian analysis (Section 6.4). Natural extensions are an own implementation of an adaptive L-stable method, positivity-preserving high-order schemes, and a pseudospectral study of non-normality.

8 References

- [1] C. F. Curtiss and J. O. Hirschfelder. Integration of stiff equations. *Proceedings of the National Academy of Sciences of the United States of America*, 38(3):235–243, 1952. DOI: [10.1073/pnas.38.3.235](https://doi.org/10.1073/pnas.38.3.235).
- [2] H. H. Robertson. The solution of a set of reaction rate equations. In J. Walsh, editor, *Numerical Analysis: An Introduction*, pages 178–182. Academic Press, London, 1966.
- [3] E. Hairer, S. P. Nørsett, and G. Wanner. *Solving Ordinary Differential Equations I: Nonstiff Problems*. 2nd revised edition, Springer-Verlag, Berlin, 1993. Springer Series in Computational Mathematics, Vol. 8. DOI: [10.1007/978-3-540-78862-1](https://doi.org/10.1007/978-3-540-78862-1).
- [4] J. C. Butcher. *Numerical Methods for Ordinary Differential Equations*. 3rd edition, John Wiley & Sons, Chichester, 2016. ISBN: [978-1-119-12150-3](https://doi.org/10.1002/9781119121534). DOI: [10.1002/9781119121534](https://doi.org/10.1002/9781119121534).
- [5] E. Hairer and G. Wanner. *Solving Ordinary Differential Equations II: Stiff and Differential-Algebraic Problems*. 2nd revised edition, Springer-Verlag, Berlin, 1996. Springer Series in Computational Mathematics, Vol. 14. DOI: [10.1007/978-3-642-05221-7](https://doi.org/10.1007/978-3-642-05221-7).
- [6] C. T. Kelley. *Solving Nonlinear Equations with Newton’s Method*. Society for Industrial and Applied Mathematics, Philadelphia, 2003. ISBN: [978-0-89871-546-0](https://doi.org/10.1137/1.9780898718898). DOI: [10.1137/1.9780898718898](https://doi.org/10.1137/1.9780898718898).
- [7] P. Deuffhard. *Newton Methods for Nonlinear Problems: Affine Invariance and Adaptive Algorithms*. Springer, Berlin, Heidelberg, 2011. Springer Series in Computational Mathematics, Vol. 35. DOI: [10.1007/978-3-642-23899-4](https://doi.org/10.1007/978-3-642-23899-4).
- [8] G. Söderlind. Automatic control and adaptive time-stepping. *Numerical Algorithms*, 31:281–310, 2002. DOI: [10.1023/A:1021160023092](https://doi.org/10.1023/A:1021160023092).
- [9] L. F. Shampine and M. W. Reichelt. The MATLAB ODE Suite. *SIAM Journal on Scientific Computing*, 18(1):1–22, 1997. DOI: [10.1137/S1064827594276424](https://doi.org/10.1137/S1064827594276424).
- [10] J. D. Lambert. *Numerical Methods for Ordinary Differential Systems: The Initial Value Problem*. John Wiley & Sons, Chichester, 1991. ISBN: [978-0-471-92990-1](https://doi.org/10.1137/0904010).
- [11] L. Petzold. Automatic selection of methods for solving stiff and nonstiff systems of ordinary differential equations. *SIAM Journal on Scientific and Statistical Computing*, 4(1):136–148, 1983. DOI: [10.1137/0904010](https://doi.org/10.1137/0904010).
- [12] A. C. Hindmarsh, P. N. Brown, K. E. Grant, S. L. Lee, R. Serban, D. E. Shumaker, and C. S. Woodward. SUNDIALS: Suite of nonlinear and differential/algebraic equation solvers. *ACM Transactions on Mathematical Software*, 31(3):363–396, 2005. DOI: [10.1145/1089014.1089020](https://doi.org/10.1145/1089014.1089020).
- [13] C. R. Harris, K. J. Millman, S. J. van der Walt, et al. Array programming with NumPy. *Nature*, 585(7825):357–362, 2020. DOI: [10.1038/s41586-020-2649-2](https://doi.org/10.1038/s41586-020-2649-2).
- [14] S. K. Lam, A. Pitrou, and S. Seibert. Numba: a LLVM-based Python JIT compiler. In *Proceedings of the Second Workshop on the LLVM Compiler Infrastructure in HPC*. Association for Computing Machinery, 2015. DOI: [10.1145/2833157.2833162](https://doi.org/10.1145/2833157.2833162).
- [15] J. D. Hunter. Matplotlib: A 2D graphics environment. *Computing in Science & Engineering*, 9(3):90–95, 2007. DOI: [10.1109/MCSE.2007.55](https://doi.org/10.1109/MCSE.2007.55).

- [16] P. Virtanen, R. Gommers, T. E. Oliphant, et al. SciPy 1.0: fundamental algorithms for scientific computing in Python. *Nature Methods*, 17:261–272, 2020. DOI: [10.1038/s41592-019-0686-2](https://doi.org/10.1038/s41592-019-0686-2).
- [17] C. W. Gear. *Numerical Initial Value Problems in Ordinary Differential Equations*. Prentice-Hall, Englewood Cliffs, NJ, 1971. ISBN: [978-0-13-626606-8](https://doi.org/10.1007/978-0-13-626606-8).
- [18] N. J. Higham. *Accuracy and Stability of Numerical Algorithms*. 2nd edition, Society for Industrial and Applied Mathematics, Philadelphia, 2002. ISBN: [978-0-89871-521-7](https://doi.org/10.1137/1.9780898718027). DOI: [10.1137/1.9780898718027](https://doi.org/10.1137/1.9780898718027).
- [19] L. A. Segel and M. Slemrod. The quasi-steady-state assumption: A case study in perturbation. *SIAM Review*, 31(3):446–477, 1989. DOI: [10.1137/1031091](https://doi.org/10.1137/1031091).
- [20] T. Turányi, A. S. Tomlin, and M. J. Pilling. On the error of the quasi-steady-state approximation. *The Journal of Physical Chemistry*, 97(1):163–172, 1993. DOI: [10.1021/j100103a028](https://doi.org/10.1021/j100103a028).
- [21] F. Verhulst. *Methods and Applications of Singular Perturbations: Boundary Layers and Multiple Timescale Dynamics*. Springer, New York, 2005. Texts in Applied Mathematics, Vol. 50. DOI: [10.1007/0-387-28313-7](https://doi.org/10.1007/0-387-28313-7).
- [22] F. Verhulst. Singular perturbation methods for slow–fast dynamics. *Nonlinear Dynamics*, 50:747–753, 2007. DOI: [10.1007/s11071-007-9236-z](https://doi.org/10.1007/s11071-007-9236-z).

9 Appendix

Appendix A: AI Transparency Log

Date/Step	AI Prompt (summary)	What we changed & why	Verification
-----------	---------------------	-----------------------	--------------

Table 10: AI Transparency Log.

Appendix B: Individual Contribution Statement (ICS)